

FINITE-REFERENCE ERROR BOUNDS FOR RECURRENT EVALUATORS IN TINY RECURSIVE MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study the fixed-snapshot error of a finite-depth recurrent value evaluator inside a discounted plan-edit MDP. For finite reference depths $m > n$, the depth- n error is bounded by its discrepancy from the depth- m evaluator plus the depth- m Bellman residual divided by $1 - \gamma^K$. A Lipschitz value head controls the discrepancy by the recurrent path length from n to m . The result permits transient computation and assumes no latent fixed point; a contractive fixed-point bound is only a specialization. Persistent latents admit the same direct residual argument on the augmented state (x, y, z) . Propagating either value bound through conservative policy improvement additionally requires exact statewise centering and pointwise deployment of the exact policy mixture. A 64-state finite-MDP calculation numerically unit-tests the composed inequalities. Empirically, our scope is limited to 4×4 Sudoku: the theory-aligned episodic configuration obtains 0% over 10 seeds, while the higher-performing persistent implementation does not establish the uniform augmented residual, exact centering, or exact-mixture deployment required by the theory. The reported UPI-TRM/PPO comparison also uses unequal environment-interaction budgets. We therefore present the theory as a finite-reference error decomposition, not as a certificate for the current implementation or as evidence of sample efficiency.

1 INTRODUCTION

Tiny Recursive Models (TRMs) (Jolicoeur-Martineau, 2025) obtain predictions through iterative latent refinement. In checker-guided tasks such as Sudoku, the same recurrent computation can serve as an internal evaluator for candidate edits to a partial plan. A finite unroll can be useful while remaining transient: there is no general reason for an arbitrary TRM to be designed around convergence to a latent fixed point. Standard approximate-policy-iteration and conservative-policy-improvement (CPI) analyses (Bertsekas & Tsitsiklis, 1996; Kakade & Langford, 2002) place evaluator error in a black-box term and therefore do not distinguish a finite recurrent path from a Bellman residual at a chosen reference depth.

We analyze Unrolled Policy Iteration with TRMs (UPI-TRM), which retains the inner recurrence, uses a scalar value head to score candidate plan edits, and places environment transitions only in the outer checker-guided MDP. The analysis fixes the MDP, recurrence, value head, initialization, and policy pair. It is not an end-to-end theorem for SGD, TD learning, BPTT, or distillation. Euclidean projection can supply forward invariance, and exact baseline summation can supply the statewise centering needed by the CPI step, but these are conditions on a frozen snapshot rather than guarantees about how training reaches it.

Our primary result (Theorem 6.1) compares $U_n(s) = V_\psi(z^{(n)}(s), x)$ with any deeper finite reference U_m for $m > n$. It bounds $\|U_n - V^\pi\|$ by $\|U_n - U_m\|$ plus the Bellman residual of U_m divided by $1 - \gamma^K$, and controls the first term by the recurrent path length for a Lipschitz head. No recurrent convergence is assumed. A forward-invariant contraction geometrically bounds that path and recovers the fixed-point formula as $m \rightarrow \infty$. Corollary 6.3 transfers the value bound to centered advantage error, and Theorem 6.5 supplies a one-step CPI statement only for exact centering and direct deployment of the exact mixture.

The evidence has two distinct roles. First, an exactly solvable finite MDP numerically unit-tests the composed value, advantage, and CPI inequalities; it is not an empirical discovery. Second,




Figure 1: TRM-to-MDP bridge used throughout the paper. Left: the original TRM’s inner latent-refinement loop nested inside its outer continuous refinement of the current answer y_t (Jolicoeur-Martineau, 2025). Right: we swap TRM’s continuous outer refinement for a discrete plan-edit loop under checker feedback, modeling it as a plan-space MDP over instance–plan states, and treat the inner loop as an internal evaluator that is unrolled for depth n . The checker maps to reward; the inner loop maps to evaluator computation inside a single state evaluation.

experiments on hard 4×4 Sudoku evaluate implementations. The theory-aligned episodic run obtains 0% success over 10 seeds. A persistent, distilled implementation obtains $57.4\% \pm 12.2\%$ in a supplied UPI–TRM reevaluation artifact. The manuscript-recorded TRM+PPO result is $32.0\% \pm 15.3\%$ at the same nominal outer-step count, but its per-seed run artifact is unavailable and the environment-interaction counts differ. We therefore claim neither equal-interaction superiority nor sample efficiency. A reported projection–clamping factorial has only aggregate provenance; we treat projection as a composite forward-pass and training-dynamics intervention and do not claim that clamping controls the global contraction modulus.

Scope. Episodic latents are Markov on (x, y) . At a fixed snapshot, persistent latents are Markov on (x, y, z) when every transition-relevant variable is included. Their direct residual certificate needs no latent fixed point or slow drift; slow drift is needed only for the stronger comparison with a memoryless fixed-point reference. For the headline run, the uniform augmented residual and exact augmented-state centering are not established, several transition details cannot be audited from the retained package, and distillation is known to violate exact-mixture deployment. The current empirical domain is only 4×4 Sudoku.

Contributions.

- A fixed-point-free finite-reference-depth value and path decomposition, with the contractive fixed-point result as a specialization (Theorem 6.1).
- A direct finite-depth certificate for persistent state on the augmented MDP, separated from an optional slow-drift comparison (Corollary B.1).
- A conditional bridge from these value errors to CPI under exact centering and exact-mixture deployment (Theorem 6.5); this is a fixed-snapshot result, not a training theorem.
- A numerical finite-MDP unit test and a 4×4 Sudoku study that reports the theory-aligned 0% result, the persistent implementation result under unequal interactions, and aggregate projection–clamping contrasts with their provenance limitations.

2 RELATED WORK

Recursive latent reasoning models. TRMs (Jolicoeur-Martineau, 2025) and Hierarchical Reasoning Models (Wang et al., 2025a) use iterative latent refinement for structured reasoning; follow-up work studies training dynamics and test-time adaptation (Qasim & Zhang, 2025; McGovern, 2025). We keep TRM’s inner latent refinement but swap its supervised outer loop for a discrete plan-edit loop

under checker feedback, and analyze how finite inner evaluation affects policy improvement. The contraction assumption is DEQ-adjacent (Bai et al., 2019) only in providing a regime with bounded truncation bias.

Policy-improvement analyses. CPI, TRPO, and safe/monotone policy improvement (Kakade & Langford, 2002; Schulman et al., 2015; Pirota et al., 2013; Metelli et al., 2021; Eshwar et al., 2025) all bound improvement under an abstract evaluation-error term, but treat that error as a single black-box quantity. The closest formal neighbor is function-approximation CPI (Metelli et al., 2021; Eshwar et al., 2025), which relates a learned V_ψ to V^π via a regression model, a sampling distribution, and an optimizer. We do something orthogonal: conditional on sup-norm residual and centering control, we expose the finite recurrent path between depths n and m . A contractive recurrence turns this path into an L_z^n specialization, but contraction is not required by the primary result. Empirical RLVR work on LLM reasoning (Wang et al., 2025b; Su et al., 2025; DeepSeek-AI, 2025) uses verifiable rewards in a similar spirit but does not analyze approximation error. A quantitative join with FA-CPI on V_ψ would need coverage and uniform-convergence assumptions we leave out.

Planning and structured-decision neighbors. Sequence-model RL ranges from architecture-level planning (Tamar et al., 2016) and inference-time search (Janner et al., 2021) to pure sequence modeling (Chen et al., 2021; Laskin et al., 2023); structured prediction and Sudoku/constraint RL (Daumé III et al., 2009; Ross et al., 2011; Chang et al., 2015; Mehta, 2021; Klein, 2022) are closer to our plan-space MDP framing. The TRM inner unroll differs from both: it is architecture-internal computation within one state evaluation.

3 BACKGROUND AND SETUP

The original TRM, given input x and plan y_t , runs an inner loop that updates a latent $z^{(k+1)} = f_\theta(z^{(k)}, y_t, x)$. After n steps it reads out a refined prediction from a task head trained on checker-labeled puzzles (Jolicoeur-Martineau, 2025). UPI-TRM keeps f_θ but replaces the task head with two pieces. The first is an optional projection Π_R that keeps latents inside a radius- R ball (Assumption 5.1, Eq. 23), so the projected update is $z^{(k+1)} = (\Pi_R \circ f_\theta)(z^{(k)}, y_t, x)$. The second is a scalar value head V_ψ that produces the unrolled evaluator output $U_n(s_t) = V_\psi(z_t^{(n)}, x)$ with $s_t = (x, y_t)$, which scores candidate plans in the outer CPI loop.

In the plan-space MDP the external state is $s_t = (x, y_t)$, the action is an edit a_t , and the transition is $y_{t+1} = \text{edit}(y_t, a_t; x)$. A checker score $c(x, y)$ supplies reward, optionally via potential shaping (Ng et al., 1999) $r = r_0 + \gamma\Phi(s') - \Phi(s)$ with $\Phi = c$. The latent z is *not* part of the environment state in our primary analysis; it is internal computation used to evaluate s_t . The resulting process is a discounted MDP $(\mathcal{S}, \mathcal{A}, P, r, \rho, \gamma)$ with $\gamma \in (0, 1)$ and $K \geq 1$. All theoretical statements fix one parameter snapshot: the MDP, recurrence, initialization, value head, current policy π , and candidate policy π_{cand} do not change while the displayed quantities are evaluated. The formal results assume finite or countable state and action spaces and use ordinary sup norms; the proofs also apply when only the relevant policy-pair closure and supported action sets are finite or countable. Rewards are bounded and the relevant closure is invariant under π almost surely. This setting covers every reported experiment. Appendix I states the additional domination and total-variation requirements for a general-kernel extension. Increasing n spends more computation inside one state evaluation but does not change the environment horizon or search depth, so the analysis tracks evaluator truncation error inside a fixed external MDP.

Episodic- z vs. persistent- z . In *episodic- z* , the latent is reinitialized from the current plan at each edit step, so (x, y) remains a Markov state. In *persistent- z* , the latent is carried across edits, matching the original TRM. With persistent latents the next-state distribution depends jointly on (x, y, z) , since the same plan paired with a different carried latent can produce a different next plan after another inner unroll. Treating (x, y) as the state in this regime is a non-Markov approximation. Appendix B lifts the finite-depth residual certificate directly to $\bar{s} = (x, y, z)$; this step needs no latent fixed point or slow drift. A separate slow-drift assumption is required only for comparison with a memoryless fixed-point reference.

Notation. For one common initialization $z^{(0)}(s) = z_{\text{init}}(s)$, let $z^{(j)}(s)$ be the latent after j applications of the frozen recurrent map and let $U_j(s) = V_\psi(z^{(j)}(s), x)$. At the absorbing state, every U_j uses the same depth-independent normalization from Appendix J. Only the op-

Algorithm 1 One UPI-TRM update on the plan-space MDP

```

1: Initialize: sample instance  $x$  and initialize plan  $y_0$ 
2: Latent source: reinitialize from the current plan (episodic- $z$ ) or carry from the previous edit (persistent- $z$ )
3: for  $t = 0, \dots, T - 1$  until solved or budget exhausted do
4:   Evaluator:  $z_t^{(n)} \leftarrow (\Pi_R \circ f_\theta)^{\circ n}(z_t, y_t, x)$ 
5:   Action:  $U_n(s_t) \leftarrow V_\psi(z_t^{(n)}, x)$ ; sample  $a_t \sim \pi_\phi(\cdot \mid y_t, z_t^{(n)}, x)$ 
6:   Transition: apply the edit, observe  $r_t$ , and store  $(s_t, a_t, r_t, s_{t+1})$ 
7: end for
8: Value update: form the  $K$ -step bootstrap target  $G^{(K)}$  and regress  $U_n$  toward it
9: Advantage estimate: compute centered  $\hat{A}$  by exact baseline summation
10: Policy update: take a policy-gradient step to obtain  $\pi_{\text{cand}}$ 
11: Conservative step:  $\pi_{\text{new}} \leftarrow (1 - \alpha)\pi_\phi + \alpha\pi_{\text{cand}}$ 
12: Optional distillation: distill  $\pi_{\text{new}}$  into a single policy network

```

tional contractive specialization uses the latent fixed point z^* and $U_*(s) = V_\psi(z^*(s), x)$. We write $\|f\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} := \sup_{s \in \mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} \cup \{s_{\text{abs}}\}} |f(s)|$ for the ordinary sup-norm over the advantage-evaluation closure used by CPI. This closure contains states reachable under π and the one-step successors needed to evaluate actions supported by π and π_{cand} , and is invariant under subsequent π dynamics almost surely (Appendix I; full notation in Appendix H).

4 UPI-TRM ALGORITHM

Algorithm 1 alternates three phases: collect edit trajectories, fit U_n against a K -step bootstrap target, and take a conservative mixture step. The formal guarantee applies to the explicit mixture; distillation is an implementation step outside the theorem.

The direct episodic and augmented-state applications of Theorem 6.5 require exact statewise baseline summation, the explicit mixture step, and no distillation. In the persistent case, centering and the mixture identity are conditional on the full augmented state. The explicit mixture (rather than a clipped surrogate, trust-region update, parameter interpolation, or distilled approximation) is required for the linear-in- α tightening; otherwise the centering defect and deployment mismatch lie outside the stated bound.

Persistent- z is handled in Appendix B, and the initialization condition in Appendix N gives $C_z \leq 2R$. Per update, the dominant cost is $O(T \cdot |A| \cdot n)$ inner-evaluator unrolls: T states per trajectory, $|A|$ unrolls each for exact baseline summation, and depth n per unroll.

5 BELLMAN RESIDUALS AND AN OPTIONAL CONTRACTIVE REGIME

We train U_n by K -step bootstrapped regression with loss $\mathcal{L}_{\text{val}} = \mathbb{E}_{s_0 \sim \mu} [(V_\psi(z^{(n)}(s_0), x_0) - \text{sg}(G^{(K)}))^2]$, where $G^{(K)} = \sum_{k < K} \gamma^k r_k + \gamma^K \bar{V}(s_K)$ is the K -step bootstrap target. The K -step Bellman operator $\mathcal{T}_K^\pi$ is the standard γ^K -contraction with fixed point V^π (Sutton & Barto, 2018; Puterman, 1994).

The primary finite-reference-depth result in Section 6 uses Bellman contraction but does not assume that the recurrent computation converges. It assumes sup-norm residual control on the relevant closure; we do not prove a general $L^2(\mu) \rightarrow \|\cdot\|_\infty$ conversion. A separate, optional assumption below makes a latent fixed-point specialization available. Bellman contraction is operator-level; inner-evaluator contraction of $z \mapsto \tilde{f}_\theta(z, y, x)$ controls a recurrent path within one frozen forward pass and by itself gives no TD- or training-convergence guarantee.

Assumption 5.1 (Forward-invariant contraction). The latent space $\mathcal{Z}$ is a complete normed space (in our implementation, $\mathbb{R}^d$ with the Euclidean norm). There exists a closed set $\mathcal{Z}_{\text{inv}} \subseteq \mathcal{Z}$ such that the projected update map $\tilde{f}_\theta(z, y, x) := (\Pi_R \circ f_\theta)(z, y, x)$ maps $\mathcal{Z}_{\text{inv}}$ to itself and is an L_z -contraction in z with $L_z \in (0, 1)$. For every (x, y) , the initialization satisfies $z_{\text{init}}(x, y) \in \mathcal{Z}_{\text{inv}}$. Let $z^*(x, y)$ denote the unique fixed point and assume

$$C_z := \sup_{(x, y)} \|\tilde{f}_\theta(z_{\text{init}}(x, y), y, x) - z_{\text{init}}(x, y)\| < \infty.$$

We enable operator-norm clamping and per-layer scaling as a contraction-oriented intervention; Π_R provides forward invariance. The retained finite-difference quantities are local proxies and do not certify the global modulus required by Assumption 5.1. We therefore distinguish the theorem’s contractive regime from the empirical *clamping enabled* condition. Value-head spectral normalization is disabled in these comparisons to avoid a separate value-head confound.

In a projection-active exploratory sweep, the local post-projection proxy $\hat{L}_z^{\text{post}}$ stays in a narrow 0.219–0.242 band, so the L_z -dependent terms are monitoring quantities rather than a tunable empirical knob. The protocol-specified episodic- z hard-suite rerun (Appendix G.3) yields a broader post-projection distribution including seeds with $\hat{L}_z^{\text{post}} \geq 1$, so the narrow-band behavior is not a general hard-suite geometry fact.

6 FINITE-DEPTH EVALUATION ERROR DECOMPOSITION AND CPI PENALTY

We connect finite evaluator depth to CPI. The CPI shell is standard (Kakade & Langford, 2002); the TRM-specific content is a finite-reference-depth error decomposition that does not require a recurrent fixed point. We focus on episodic- z and defer persistent- z to Appendix B.

6.1 VALUE ERROR BOUNDS (EPISODIC- z)

K -step Bellman contraction gives, for any bounded measurable V on the forward-invariant closure,

$$\|V - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq \frac{1}{1 - \gamma^K} \|V - \mathcal{T}_K^\pi V\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}, \quad (1)$$

and therefore permits any finite unroll depth to serve as the reference.

Theorem 6.1 (Finite-reference-depth value decomposition). *Fix a finite or countable MDP, or a finite or countable policy-pair closure, and one parameter snapshot, policies (π, π_{cand}) , recurrent map, initialization convention, $K \geq 1$, and a latent norm on $\mathcal{Z} \subseteq \mathbb{R}^d$. Let $V_\psi : \mathcal{Z} \times \mathcal{X} \rightarrow \mathbb{R}$ be the scalar value head. For any finite integers $0 \leq n < m$, let $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} \cup \{s_{\text{abs}}\}$ contain the one-step successors required by both policies and be invariant under subsequent π transitions almost surely. Assume rewards are bounded there; $U_j(s) = V_\psi(z^{(j)}(s), x)$ is bounded there for $j \in \{n, \dots, m\}$; and $U_j(s_{\text{abs}}) = -C_{\text{max}}$ for every depth. All displayed norms are ordinary sup norms. Then*

$$\|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq \underbrace{\|U_n - U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}}_{\text{finite-depth discrepancy}} + \underbrace{\frac{\|U_m - \mathcal{T}_K^\pi U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}}{1 - \gamma^K}}_{\text{reference-depth residual}}. \quad (2)$$

If, in addition, the scalar head is L_V -Lipschitz in its latent argument uniformly in x , then

$$\|U_n - U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq L_V \sum_{j=n}^{m-1} \sup_{s \in \mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}} \|z^{(j+1)}(s) - z^{(j)}(s)\|. \quad (3)$$

The theorem is a residual bound plus a telescoping path bound; we make no novelty claim for the triangle inequality itself. Its role is to expose a certifiable finite reference m : the recurrence may perform useful transient computation and need not converge before, at, or after m . The same frozen policy π appears in V^π and $\mathcal{T}_K^\pi$, and the latter bootstraps from U_m itself.

For the optional contractive specialization, additionally assume the fixed-point map is measurable and U_* is bounded on the closure (the latter follows from Assumption A.3); set $U_*(s_{\text{abs}}) = -C_{\text{max}}$ and define

$$\varepsilon_{\text{res}}^* \triangleq \|U_* - \mathcal{T}_K^\pi U_*\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}. \quad (4)$$

Under Assumption 5.1 and the Lipschitz-head condition, consecutive latent differences satisfy $\|z^{(j+1)} - z^{(j)}\| \leq L_z^j C_z$ uniformly. Hence Theorem 6.1 gives the finite- m specialization

$$\|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq L_V C_z \frac{L_z^n (1 - L_z^{m-n})}{1 - L_z} + \frac{\|U_m - \mathcal{T}_K^\pi U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}}{1 - \gamma^K}. \quad (5)$$

Uniform convergence $U_m \rightarrow U_*$ and Bellman contraction imply convergence of the residual norm. Taking $m \rightarrow \infty$ recovers

$$\|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq \underbrace{\frac{\varepsilon_{\text{res}}^*}{1 - \gamma^K}}_{\text{fixed-point residual}} + \underbrace{L_V \frac{L_z^n}{1 - L_z} C_z}_{\text{geometric path bound}}. \quad (6)$$

Thus the earlier fixed-point decomposition is a specialization obtained by geometrically bounding the recurrent path and sending the finite reference depth to infinity. It is not an assertion that arbitrary TRMs are designed to converge to a fixed point. Under $\|z_{\text{init}}\|_2 \leq R$ with $\mathcal{Z}_{\text{inv}} = B_R$, projection gives $C_z \leq 2R$ (Remark M.2).

A projection-active specialization makes the role of R explicit. When projection is active (that is, $\|f_\theta(z, y, x)\|_2 \geq \rho_R > R$ for all $z \in B_R$ and relevant (x, y)), Π_R also contracts the evaluator. Proposition M.3 gives

$$L_z \leq L_z^{\text{sat}}(R) := \frac{R}{\rho_R} L_z^{\text{pre}}(R), \quad (7)$$

and Proposition M.4 bounds the residual relative to the boundary-normalized zero-latent baseline

$$B_0(x, y) := V_\psi(0, x), \quad B_0(s_{\text{abs}}) := -C_{\text{max}} :$$

$$\varepsilon_{\text{res}}^* \leq \varepsilon_0 + (1 + \gamma^K) L_V(R) R, \quad (8)$$

with $\varepsilon_0 := \|B_0 - \mathcal{T}_K^\pi B_0\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}$. Projection thus enters both the architectural term (through $L_V(R)R$) and the truncation term (through $L_z^{\text{sat}}(R)$ and $C_z \leq 2R$). The experiments report finite-sample proxies for this regime. They do not certify the uniform ρ_R and Lipschitz assumptions; the held-out U_n Bellman residual upper-bounds $\varepsilon_{\text{res}}^*$ only up to the truncation gap (Remark K.3).

Corollary 6.2 (Finite-depth residual certificate). *Fix a policy π and a finite-depth evaluator U_n on the advantage-evaluation closure $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$. Define*

$$\varepsilon_{\text{res}, n} \triangleq \|U_n - \mathcal{T}_K^\pi U_n\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}. \quad (9)$$

Then

$$\|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq \frac{\varepsilon_{\text{res}, n}}{1 - \gamma^K}. \quad (10)$$

If $\hat{A} = \hat{A}_{U_n}$ is the exact one-step centered estimator of Eq. (67), then Assumption 6.4 holds with

$$\varepsilon_A \triangleq \frac{2\gamma \varepsilon_{\text{res}, n}}{1 - \gamma^K}. \quad (11)$$

Consequently, under the hypotheses of Theorem 6.5,

$$\eta(\pi_{\text{new}}) \geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{2\alpha\gamma}{1 - \gamma} \frac{\varepsilon_{\text{res}, n}}{1 - \gamma^K} - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2. \quad (12)$$

This certificate follows from Bellman contraction alone and applies in finite MDPs or deterministic edit environments where the K -step backup is computable. It does not require the Lipschitz or contractive architectural assumptions. A finite-MDP unit test checks 80 value/decomposition rows and 21 CPI rows with zero violations (Appendix C.1).

The CPI guarantees below are deterministic conditional statements that need sup-norm residual control on $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$. In finite or exactly evaluable MDPs Corollary 6.2 can certify it; otherwise it remains an assumption. Because the closure depends on both π and π_{cand} , residual control is a property of the policy pair. The certificate is post-candidate: once π_{cand} is chosen, the relevant closure can be checked; future candidates require their own check.

Corollary 6.3 (Advantage-error bridge). *Under the hypotheses of Theorem 6.1, define*

$$B_{n, m} := \|U_n - U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} + \frac{\|U_m - \mathcal{T}_K^\pi U_m\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}}{1 - \gamma^K}. \quad (13)$$

Let $\hat{A} = \hat{A}_{U_n}$ be the exact one-step centered estimator of Eq. (67), with the transition expectation evaluated exactly (or applied to a deterministic transition) and the baseline computed exactly under π . Then Assumption 6.4 holds with the valid choice

$$\varepsilon_A = 2\gamma B_{n,m}. \quad (14)$$

If Assumption 5.1 and the L_V -Lipschitz head condition also hold, sending $m \rightarrow \infty$ gives the fixed-point specialization

$$\varepsilon_A \leq \underbrace{\frac{2\gamma \varepsilon_{\text{res}}^*}{1 - \gamma^K}}_{\text{architectural}} + 2\gamma L_V \underbrace{\frac{L_z^n}{1 - L_z}}_{\text{unrolling}} C_z. \quad (15)$$

Here the value-error norm is taken over $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ (Definition I.2). For a centered K -step estimator, 2γ becomes $2\gamma^K$; GAE-style constants can likewise be absorbed into ε_A (Remark L.3). Combining the finite-reference choice with Theorem 6.5 gives the fixed-point-free penalty

$$\eta(\pi_{\text{new}}) \geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{2\alpha\gamma}{1 - \gamma} B_{n,m} - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2. \quad (16)$$

Thus a finite recurrent path and a residual at any deeper finite depth feed directly into CPI. Geometric L_z^n dependence is only the contractive specialization.

6.2 CONSERVATIVE POLICY IMPROVEMENT

For $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$, standard CPI gives monotonic improvement up to an $O(\alpha^2)$ penalty. The decompositions above enter through centered evaluation error: under exact centering, candidate-policy advantage error contributes the explicit linear term $\alpha \varepsilon_{A, \text{cand}} / (1 - \gamma)$.

Assumption 6.4 (Uniform advantage error on states visited by π). For some $\varepsilon_A \geq 0$:

$$|\hat{A}(s, a) - A^\pi(s, a)| \leq \varepsilon_A \quad (17)$$

for all (s, a) with $s \in \mathcal{R}_\pi$ and $a \in \text{supp}(\pi_{\text{cand}}(\cdot|s)) \cup \text{supp}(\pi(\cdot|s))$.

The relevant support depends on π_{cand} (standard CPI circularity; see Appendix I). Write $\hat{L}_\pi(\pi') := \eta(\pi) + \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_\pi, a \sim \pi'(\cdot|s)} [\hat{A}(s, a)]$ for the CPI surrogate built from the approximate advantage.

Theorem 6.5 (CPI with centered evaluation error). *Let $\alpha \in [0, 1]$ and let $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$ be the exact mixture policy, with no distillation error. Suppose Assumption 6.4 holds. Assume the advantage estimator is statewise-centered under π :*

$$\mathbb{E}_{a \sim \pi(\cdot|s)} [\hat{A}(s, a)] = 0 \quad \text{for all } s \in \mathcal{R}_\pi. \quad (18)$$

We define the candidate-policy bias in advantage estimation

$$\varepsilon_{A, \text{cand}} \triangleq \sup_{s \in \mathcal{R}_\pi} \left| \mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)} [\hat{A}(s, a) - A^\pi(s, a)] \right|. \quad (19)$$

Then:

$$\eta(\pi_{\text{new}}) \geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{\alpha}{1 - \gamma} \varepsilon_{A, \text{cand}} - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2, \quad (20)$$

where $\varepsilon_{\text{CPI}} \triangleq \sup_s |\mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)} A^\pi(s, a)|$.

The linear-in- α contribution $\alpha \varepsilon_{A, \text{cand}} / (1 - \gamma)$, instead of the naive $\varepsilon_A / (1 - \gamma)$, requires the exact-centering condition (18). This condition holds whenever the estimator is constructed to be exactly statewise centered. For Eq. (67), the baseline expectation must be evaluated exactly; in our finite action space we use explicit summation. Sampling-based or learned approximations instead incur the centering defect of Corollary K.2.

The constants split into estimable and assumed (Table 3): $\hat{L}_z^{\text{post}}$, $\hat{\rho}_R$, $\hat{C}_z$, $\hat{L}_V$, and $\hat{\varepsilon}_{\text{cent}}$ admit finite-difference or batch proxies. The architectural residual $\varepsilon_{\text{res}}^*$ is a sup-norm over the advantage closure that we do not estimate uniformly; the held-out finite-batch residual upper-bounds it only up to the truncation gap (Remark K.3). On the theory-aligned episodic- z hard-suite rerun, the predicted CPI penalty is non-finite on 2/10 seeds at every α (Table 20). At this scale, the bound functions as a structural decomposition; its constants are not calibrated enough to predict Sudoku-scale improvement.

Corollary 6.6 (Sufficient condition for one-step improvement). *Under Theorem 6.5, if*

$$\hat{L}_\pi(\pi_{\text{new}}) - \eta(\pi) \geq \frac{\alpha}{1-\gamma} \varepsilon_{A,\text{cand}} + \frac{2\varepsilon_{\text{CPI}}\gamma}{(1-\gamma)^2} \alpha^2, \quad (21)$$

then $\eta(\pi_{\text{new}}) \geq \eta(\pi)$. In particular, under Corollary 6.3, it suffices to replace $\varepsilon_{A,\text{cand}}$ in (21) by $2\gamma B_{n,m}$. Under the optional contractive specialization, one may further upper-bound this quantity by

$$2\gamma \left(\frac{\varepsilon_{\text{res}}^*}{1-\gamma^K} + L_V \frac{L_z^n}{1-L_z} C_z \right).$$

Theorem 6.5 applies to the exact mixture; distillation error lies outside it. Exact baseline summation over a discrete action space satisfies (18); approximate centering replaces the linear term by $[(1-\alpha)\varepsilon_{\text{cent}} + \alpha\varepsilon_{A,\text{cand}}]/(1-\gamma)$ via Corollary K.2. The theorem is most informative in the small- α regime, where the linear term dominates the α^2 penalty.

Corollary 6.7 (One-step parameter dependence). *Under Corollary 6.3 and Theorem 6.5,*

$$\eta(\pi_{\text{new}}) \geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{2\alpha\gamma}{1-\gamma} \left(\frac{\varepsilon_{\text{res}}^*}{1-\gamma^K} + L_V \frac{L_z^n}{1-L_z} C_z \right) - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1-\gamma)^2} \alpha^2. \quad (22)$$

Corollary 6.7 is the optional fixed-point specialization of the finite-reference penalty $2\alpha\gamma B_{n,m}/(1-\gamma)$. It makes the one-step dependence on (L_z, n, K, α) explicit and reduces to standard CPI in the limit $L_z, \varepsilon_{\text{res}}^* \rightarrow 0$. Two constants in (22) are observed rather than tuned: L_V is only locally estimated by finite differences, and the current clamping mechanism produces little range in one projection-active proxy $\hat{L}_z^{\text{post}}$ (≈ 0.22 – 0.24). On the hard suite, we therefore treat the L_z -dependent terms as monitoring quantities.

A projection-aware version of (22) substitutes $\varepsilon_0 + (1+\gamma^K)L_V(R)R$ for $\varepsilon_{\text{res}}^*$ and $L_z^{\text{sat}}(R)$ for L_z ; the full statement is Eq. (70) in Appendix M.5. In the saturated regime ($\rho_R \gg R$), R/ρ_R sets $L_z^{\text{sat}}(R)$ more directly than the clamping target does. Smaller L_z or larger n reduces L_z^n ; smaller α is safer; larger K tightens $1/(1-\gamma^K)$ at the cost of a K -dependent residual. Persistent- z uses the direct augmented-state residual of Appendix B; slow drift enters only the optional comparison to a memoryless fixed-point reference.

7 EXPERIMENTS

All task results in this paper are on 4×4 Sudoku. The decision-relevant evidence is a theory-aligned episodic run at 0%, a persistent distilled run at 57.4% under a different interaction budget, and manuscript-reported aggregate projection $\times$ clamping cells. These implementations do not validate the fixed-snapshot theorem; Appendix Table 2 maps each claim to its conditions, provenance, and status.

Setup. The persistent result, reported factorial, and nominal baselines use no-mask 6–8-empty Sudoku ($T=16$) for 20k outer steps; the episodic run stops at 20k environment steps. Nominal interactions are unequal. Architecture-matched baselines share the TRM backbone but not the update rule; external SB3 baselines are implementation checks. The clamping-enabled condition uses an operator-norm target of 0.9; this label does not assert that the global modulus is below one.

7.1 PERSISTENT HARD-SUITE IMPLEMENTATION AND LIMITATIONS

The persistent, distilled implementation produces the higher task result in Table 1. The no-mask 6–8-empty protocol is the first setting in this paper where masked random play scores 0%.

At the nominal endpoint, the ten manuscript-recorded seed values imply a +25.4pp UPI-TRM/TRM+PPO gap (Welch $t \approx 4.10$, $p \approx 7.3 \times 10^{-4}$, 95% CI [12.4, 38.5]pp). The arithmetic is reproducible, but the PPO per-seed run artifact is unavailable. This joint, unequal-interaction implementation contrast isolates no component and implies no efficiency advantage (Appendix G.2).

For the persistent run, clock/carry details are unaudited, the uniform augmented residual and centering are unmeasured, and distillation violates exact-mixture deployment. Thus 57.4% is an implementation result, not verification of Corollary B.1 or Theorem 6.5; exact centering alone would not repair the mixture mismatch.

Table 1: No-mask hard 4×4 results on 6–8 empties ($T=16$, mean $\pm$ std across seeds, nominal 20k outer-step protocol; persistent- z). A per-seed reevaluation artifact supports the UPI–TRM row. Baseline rows are manuscript-recorded aggregates whose per-seed run artifacts are unavailable. The protocol does not equalize environment interactions across rows.

Method	Outer Steps	Seeds	Success (%)	Return	Invalid
UPI–TRM	20k	10	57.4 ± 12.2	-1.82 ± 2.66	0.000 ± 0.000
<i>Architecture-matched, manuscript-recorded aggregates (same TRM backbone):</i>					
TRM+PPO	20k	10	32.0 ± 15.3	-7.87 ± 3.65	0.000 ± 0.000
TRM+A2C	20k	10	0.0 ± 0.0	-16.81 ± 0.29	0.000 ± 0.000
TRM+DQN	20k	10	0.0 ± 0.0	-18.38 ± 0.63	0.000 ± 0.000
<i>External (SB3, flat MLP policy; high invalid-action rates):</i>					
SB3 PPO/A2C/DQN/DQN($n=5$)	20k	10	0.0 ± 0.0	≈ -17.9	$0.78\text{--}1.00$

7.2 PROJECTION AND CLAMPING ABLATION

The manuscript reports aggregate cells from a 2×2 projection $\times$ clamping factorial (Appendix Table 17); the per-seed artifact and test script are unavailable. Projection’s composite $+13.0\text{pp}$ contrast ($36.2\% \rightarrow 49.2\%$) mixes clipping and training dynamics. The clamping-enabled contrast is $+2.2\text{pp}$ and inconclusive; the source-reported $p \approx 0.6$ is not independently auditable. The theorem’s contraction condition remains a conditional error-control mechanism, not a demonstrated performance dial.

7.3 THEORY-ALIGNED EPISODIC NEGATIVE RESULT

The episodic run uses reinitialization, exact baselines/targets, and the direct mixture, obtaining 0.0% at 20k environment steps over seeds 41–50. Its four joint changes and different budget preclude a controlled comparison, and its residual is not the persistent augmented residual (Appendix G.3).

7.4 INTERACTION-BUDGET AUDIT

PPO/A2C/DQN collect 64/32/4 transitions per nominal step. UPI–TRM is absent from the common-budget reruns, whose provenance is missing; there is no matched UPI–TRM–PPO result (Appendix G.1).

7.5 FINITE-MDP NUMERICAL UNIT TEST

An exact 64-state, 4-action unit test has zero violations in 80 value and 21 selected-slice CPI rows; it is not an empirical discovery or Cartesian sweep (Appendix C.1).

8 DISCUSSION, LIMITATIONS, AND CONCLUSION

The main result is a fixed-snapshot decomposition: a finite-depth recurrent evaluator can be compared with any deeper finite reference through a recurrent-path term and a reference-depth Bellman residual. Contraction supplies one optional geometric specialization rather than a model of arbitrary TRM computation. Persistent latents admit the same direct residual argument on a clock-complete augmented MDP, but the headline implementation does not provide the uniform augmented residual, exact centering, or exact-mixture deployment required to instantiate the CPI bound.

The present empirical evidence is limited to 4×4 Sudoku. The theory-aligned episodic configuration fails under the tested budget, the UPI–TRM/PPO comparison is not interaction matched, projection has a composite forward-pass and training-dynamics effect, and the clamping-enabled contrast is inconclusive and lacks per-seed provenance. The two decision-relevant next tests are an interaction-matched comparison and a one-factor-at-a-time bridge between the persistent and episodic implementations.

REFERENCES

- Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, volume 32, pp. 690–701, 2019.
- Dimitri P. Bertsekas and John N. Tsitsiklis. *Neuro-Dynamic Programming*. Athena Scientific, 1996. ISBN 1886529108.
- Kai-Wei Chang, Akshay Krishnamurthy, Alekh Agarwal, Hal Daumé III, and John Langford. Learning to search better than your teacher. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pp. 2058–2066. PMLR, 2015.
- Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Michael Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling. In *Advances in Neural Information Processing Systems*, volume 34, pp. 15084–15097, 2021.
- Hal Daumé III, John Langford, and Daniel Marcu. Search-based structured prediction. *Machine Learning*, 75(3):297–325, 2009.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning, 2025.
- S. R. Eshwar, Gagan Thoppe, Ananyabrata Barua, Aditya Gopalan, and Gal Dalal. Monotone and conservative policy iteration beyond the tabular case, 2025.
- Michael Janner, Qiyang Li, and Sergey Levine. Offline reinforcement learning as one big sequence modeling problem. In *Advances in Neural Information Processing Systems*, volume 34, pp. 1273–1286, 2021.
- Alexia Jolicoeur-Martineau. Less is more: Recursive reasoning with Tiny networks. *arXiv preprint arXiv:2510.04871*, 2025.
- Sham Kakade and John Langford. Approximately optimal approximate reinforcement learning. In *Proceedings of the 19th International Conference on Machine Learning*, pp. 267–274, 2002.
- Kevin Klein. Sudoku #3: Poor man’s RL. https://kevinkle.in/posts/2022-01-09-sudoku_rl/, 2022. Blog post. Accessed 2026-01-09.
- Michael Laskin, Luyu Wang, Junhyuk Oh, Emilio Parisotto, Stephen Spencer, Richie Steigerwald, DJ Strouse, Steven Stenberg Hansen, Angelos Filos, Ethan Brooks, Maxime Gazeau, Himanshu Sahni, Satinder Singh, and Volodymyr Mnih. In-context reinforcement learning with algorithm distillation. In *International Conference on Learning Representations*, 2023.
- Ronan Killian McGovern. Test-time adaptation of tiny recursive models. *arXiv preprint arXiv:2511.02886*, 2025.
- Anav Mehta. Reinforcement learning for constraint satisfaction game agents (15-puzzle, minesweeper, 2048, and sudoku). *arXiv preprint arXiv:2102.06019*, 2021. URL <https://arxiv.org/abs/2102.06019>.
- Alberto Maria Metelli, Matteo Pirota, Daniele Calandriello, and Marcello Restelli. Safe policy iteration: A monotonically improving approximate policy iteration approach. *Journal of Machine Learning Research*, 22(97):1–83, 2021.
- Takeru Miyato, Toshiki Kataoka, Masanori Koyama, and Yuichi Yoshida. Spectral normalization for generative adversarial networks. In *International Conference on Learning Representations*, 2018.
- Andrew Y. Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the Sixteenth International Conference on Machine Learning*, pp. 278–287. Morgan Kaufmann, 1999.
- Matteo Pirota, Marcello Restelli, Alessio Pecorino, and Daniele Calandriello. Safe policy iteration. In *Proceedings of the 30th International Conference on Machine Learning*, volume 28 of *Proceedings of Machine Learning Research*, pp. 307–315. PMLR, 2013.

540 Martin L. Puterman. *Markov Decision Processes: Discrete Stochastic Dynamic Programming*.
541 Wiley-Interscience, 1994. ISBN 0471619779.

542

543 Kaleem Ullah Qasim and Jiashu Zhang. Accelerating training speed of tiny recursive models with
544 curriculum guided adaptive recursion. *arXiv preprint arXiv:2511.08653*, 2025.

545

546 Stephane Ross, Geoffrey J. Gordon, and J. Andrew Bagnell. A reduction of imitation learning and
547 structured prediction to no-regret online learning. In *Proceedings of the Fourteenth International
548 Conference on Artificial Intelligence and Statistics*, volume 15 of *Proceedings of Machine Learning
549 Research*, pp. 627–635. PMLR, 2011.

550

551 John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel. Trust region
552 policy optimization. In *Proceedings of the 32nd International Conference on Machine Learning*,
553 pp. 1889–1897, 2015.

554

555 Yi Su, Dian Yu, Linfeng Song, Juntao Li, Haitao Mi, Zhaopeng Tu, Min Zhang, and Dong Yu.
556 Crossing the reward bridge: Expanding rl with verifiable rewards across diverse domains, 2025.

557

558 Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2
559 edition, 2018.

560

561 Aviv Tamar, Yi Wu, Garrett Thomas, Sergey Levine, and Pieter Abbeel. Value iteration networks. In
562 *Advances in Neural Information Processing Systems*, volume 29, 2016.

563

564 Guan Wang, Jin Li, Yuhao Sun, Xing Chen, Changling Liu, Yue Wu, Meng Lu, Sen Song, and Yasin
565 Abbasi-Yadkori. Hierarchical Reasoning Model. *arXiv preprint arXiv:2506.21734*, 2025a.

566

567 Yiping Wang, Qing Yang, Zhiyuan Zeng, Liliang Ren, Liyuan Liu, Baolin Peng, Hao Cheng, Xuehai
568 He, Kuan Wang, Jianfeng Gao, Weizhu Chen, Shuohang Wang, Simon Shaolei Du, and Yelong
569 Shen. Reinforcement learning for reasoning in large language models with one training example,
570 2025b.

571

572

573

574

575

576

577

578

579

580

581

582

583

584

585

586

587

588

589

590

591

592

593

A SUPPLEMENTARY MATERIAL

Assumption A.1 (Bounded rewards). Potential-shaped rewards are uniformly bounded: $|r(s, a, s')| \leq R_{\max}$ for all transitions.

$$\begin{aligned} \Pi_R(z) &= \arg \min_{u: \|u\|_2 \leq R} \|u - z\|_2 \\ &= \begin{cases} z, & \|z\|_2 \leq R, \\ R \frac{z}{\|z\|_2}, & \|z\|_2 > R. \end{cases} \end{aligned} \quad (23)$$

$$G^{(K)}(s_0) = \sum_{k=0}^{K-1} \gamma^k r_k + \gamma^K \bar{V}(s_K). \quad (24)$$

$$C_z := \sup_{(x,y)} \|\tilde{f}_\theta(z_{\text{init}}(x, y), y, x) - z_{\text{init}}(x, y)\| < \infty. \quad (25)$$

Proposition A.2 (Fixed point and finite-unrolling bias (Banach)). *Under Assumption 5.1, the map $T_{x,y}(z) = \tilde{f}_\theta(z, y, x)$ has a unique fixed point $z^*(x, y)$, with linear convergence*

$$\|z^{(t)} - z^*(x, y)\| \leq L_z^t \|z^{(0)} - z^*(x, y)\|. \quad (26)$$

Moreover, for $z^{(0)} = z_{\text{init}}(x, y)$,

$$\|z^{(n)} - z^*(x, y)\| \leq \frac{L_z^n}{1 - L_z} C_z. \quad (27)$$

$$|V_\psi(z, x) - V_\psi(z', x)| \leq L_V \|z - z'\|. \quad (28)$$

Assumption A.3 (Bounded value head on the invariant region). There exists $V_{\max} < \infty$ such that $|V_\psi(z, x)| \leq V_{\max}$ and $|V_{\bar{\psi}}(z, x)| \leq V_{\max}$ on the invariant region.

For $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$:

$$\eta(\pi_{\text{new}}) \geq L_\pi(\pi_{\text{new}}) - \frac{2\varepsilon_{\text{CPI}} \gamma}{(1 - \gamma)^2} \alpha^2, \quad (29)$$

where $\varepsilon_{\text{CPI}} = \sup_s |\mathbb{E}_{a \sim \pi_{\text{cand}}} A^\pi(s, a)|$.

$$\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}, \quad \alpha \in [0, 1]. \quad (30)$$

B PERSISTENT- z ON THE AUGMENTED MDP

B.1 DIRECT FINITE-DEPTH RESIDUAL CERTIFICATE

The formal statement below assumes that the reachable augmented policy-pair closure is finite or countable, as in Section 3. A continuous-kernel extension requires Remark I.3.

Fix $n, (\theta, \psi)$, the policy parameters, and all environment kernels. Any timeout or remaining edit budget that affects termination is included in y , and edit updates that clock; we suppress it in the notation. Define

$$h_n(x, y, z) := T_{x,y}^{\text{on}}(z) = (\Pi_R \circ f_\theta)^{\text{on}}(z, y, x), \quad \bar{s} = (x, y, z), \quad (31)$$

and the augmented transition by first computing h_n , sampling an action, and then setting

$$\bar{s}' = (x, \text{edit}(y, a; x), h_n(x, y, z)).$$

The current and candidate action kernels share this same frozen h_n and hence the same augmented MDP:

$$\bar{\pi}(a \mid \bar{s}) = \pi_\phi(a \mid y, h_n(x, y, z), x), \quad \bar{\pi}_{\text{cand}}(a \mid \bar{s}) = \pi_{\text{cand}}(a \mid y, h_n(x, y, z), x).$$

Changing the recurrent map between these two policies would change the transition kernel and is outside ordinary CPI. The initial law is

$$\bar{\rho}(dx, dy, dz) = \rho(dx, dy) \delta_{z_{\text{init}}(x, y)}(dz), \quad (32)$$

or, for stochastic initialization, $\bar{\rho}(dx, dy, dz) = \rho(dx, dy) \nu_0(dz \mid x, y)$ with the recorded kernel ν_0 .

Assume the edit/reward kernels, h_n , V_ψ , and both policy kernels are measurable; rewards and the evaluator are bounded; and the augmented advantage-evaluation closure is measurable, contains the one-step successors under both action supports, and is $\bar{\pi}$ -forward invariant almost surely. Define

$$\tilde{U}_n(\bar{s}) := V_\psi(h_n(x, y, z), x), \quad \tilde{U}_n(\bar{s}_{\text{abs}}) := -C_{\text{max}}, \quad (33)$$

and let $\bar{T}_K^{\bar{\pi}}$ denote the K -step Bellman operator of this frozen augmented MDP. Its direct residual is

$$\bar{\varepsilon}_{\text{res}, n} := \|\tilde{U}_n - \bar{T}_K^{\bar{\pi}} \tilde{U}_n\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}}. \quad (34)$$

Here $\eta(\bar{\pi}')$ denotes the discounted return of any augmented policy $\bar{\pi}'$ from $\bar{\rho}$, and $d_{\bar{\pi}}$ is the discounted occupancy of the current policy. For an augmented advantage estimator $\hat{A}$, define

$$\hat{L}_{\bar{\pi}}(\bar{\pi}') := \eta(\bar{\pi}) + \frac{1}{1 - \gamma} \mathbb{E}_{\bar{s} \sim d_{\bar{\pi}}, a \sim \bar{\pi}'(\cdot \mid \bar{s})} [\hat{A}(\bar{s}, a)], \quad (35)$$

$$\bar{\varepsilon}_{\text{CPI}} := \sup_{\bar{s}} \left| \mathbb{E}_{a \sim \bar{\pi}_{\text{cand}}(\cdot \mid \bar{s})} \hat{A}^{\bar{\pi}}(\bar{s}, a) \right|. \quad (36)$$

Corollary B.1 (Direct persistent- z certificate on the augmented MDP). *Under the preceding frozen-MDP, measurability, boundedness, and closure assumptions,*

$$\|\tilde{U}_n - V^{\bar{\pi}}\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}} \leq \frac{\bar{\varepsilon}_{\text{res}, n}}{1 - \gamma^K}. \quad (37)$$

If $\hat{A}_{\tilde{U}_n}$ uses the exact transition expectation and exact baseline and is therefore centered pointwise under $\bar{\pi}$ on the full augmented state, then Assumption 6.4 holds on the augmented closure with the valid choice

$$\bar{\varepsilon}_A := \frac{2\gamma \bar{\varepsilon}_{\text{res}, n}}{1 - \gamma^K}. \quad (38)$$

If, additionally, the deployed policy is the pointwise exact mixture $\bar{\pi}_{\text{new}} = (1 - \alpha)\bar{\pi} + \alpha\bar{\pi}_{\text{cand}}$ with no parameter interpolation or distillation, Theorem 6.5 gives

$$\eta(\bar{\pi}_{\text{new}}) \geq \hat{L}_{\bar{\pi}}(\bar{\pi}_{\text{new}}) - \frac{2\alpha\gamma}{1 - \gamma} \frac{\bar{\varepsilon}_{\text{res}, n}}{1 - \gamma^K} - \frac{2\bar{\varepsilon}_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2. \quad (39)$$

This direct result is Bellman contraction on the actual finite-depth augmented evaluator. It requires no latent fixed point, recurrent contraction, Lipschitz head, projection, or slow-drift assumption. The augmented residual differs from the episodic residual on (x, y) because both the policy and transition depend on the carried latent. We do not measure $\bar{\varepsilon}_{\text{res}, n}$ uniformly for the headline run. Exact augmented-state centering is also not established there, and the headline implementation uses distillation, so the exact-mixture condition is known not to hold.

Remark B.2 (Scope of the direct augmented-state certificate). Corollary B.1 is a result on $\bar{s} = (x, y, z)$, not on the plan marginal (x, y) . Its exact-centering condition averages actions conditional on the full $\bar{s}$, and its mixture identity must hold pointwise on the augmented closure. Finite-batch maxima are diagnostics, not uniform certificates. The retained implementation package does not contain the training repository needed to verify latent-carry order, clock encoding, or whether candidate construction freezes the recurrent map; these remain implementation-audit requirements rather than assumed facts.

B.2 OPTIONAL SLOW-DRIFT COMPARISON TO A FIXED-POINT REFERENCE

Slow drift is needed only for the stronger architectural comparison between the actual transient evaluator and the memoryless reference $\tilde{U}_*(x, y, z) := V_\psi(z^*(x, y), x)$, with $\tilde{U}_*(\bar{s}_{\text{abs}}) = -C_{\text{max}}$. For this subsection, additionally assume $z^*(x, y)$ is measurable, $\tilde{U}_*$ is bounded and measurable on the augmented closure, and the tracking conditions below hold uniformly over trajectories formed by one action supported by either policy followed by $\bar{\pi}$ dynamics. These are the one-deviation trajectories generating the augmented advantage closure. Let $e : \mathcal{Y} \rightarrow \mathcal{E}_y$ be a fixed plan embedding and write $\|y - y'\|_{\mathcal{E}} := \|e(y) - e(y')\|_{\mathcal{E}_y}$.

Assumption B.3 (Slow drift of fixed points). Uniformly on the augmented policy-pair closure, there exist $L_{z^*} \geq 0$, $\Delta y_{\text{max}} \geq 0$, and $E_0 < \infty$ such that

1. $\|z^*(x, y) - z^*(x, y')\| \leq L_{z^*} \|y - y'\|_{\mathcal{E}}$;
2. every admissible edit satisfies $\|y_{t+1} - y_t\|_{\mathcal{E}} \leq \Delta y_{\text{max}}$; and
3. the initial carried latent satisfies $\|z_0 - z^*(x, y_0)\| \leq E_0$.

Remark B.4 (A conditional sufficient bound for L_{z^*}). If y has a specified differentiable continuous representation, f_θ is C^1 near an interior fixed point, and $\|\partial_z f_\theta\|_{\text{op}} \leq L_z < 1$, the implicit-function calculation gives

$$L_{z^*} \leq \frac{\|\partial_y f_\theta\|_{\text{op}}}{1 - L_z}. \quad (40)$$

For discrete Sudoku plans this calculation additionally requires a differentiable extension off the discrete embedding; the manuscript does not provide or verify such an extension.

Remark B.5 (Discrete-plan scope). For a discrete domain, Assumption B.3 is simply an additional regularity assumption in the stated embedding. Successful training does not by itself establish it. We do not verify it for the headline Sudoku run.

Remark B.6 (Scope of the slow-drift assumption). Only the augmented closure for the frozen policy pair is relevant; no global smoothness claim outside that closure is needed.

Define $e_t = \|z_t - z^*(x, y_t)\|$ and $\kappa_n = L_{z^*}$.

Lemma B.7 (Drift bound for persistent- z). Under Assumptions 5.1 and B.3,

$$e_{t+1} \leq \kappa_n e_t + L_{z^*} \Delta y_{\text{max}}, \quad (41)$$

and

$$\sup_t \|z_t^{(n)} - z^*(x, y_t)\| \leq C_{\text{drift}}(n) := \kappa_n \max\left(E_0, \frac{L_{z^*} \Delta y_{\text{max}}}{1 - \kappa_n}\right) \leq \kappa_n E_0 + \frac{\kappa_n L_{z^*} \Delta y_{\text{max}}}{1 - \kappa_n}. \quad (42)$$

Remark B.8 (When the drift bound is informative). The bound is useful only when its assumed constants are finite and calibrated on the policy-pair closure. Otherwise the direct residual certificate remains mathematically valid but uncalibrated, while the fixed-point comparison is vacuous.

Define

$$\bar{\varepsilon}_{\text{res},*} := \|\tilde{U}_* - \bar{T}_K^{\bar{\pi}} \tilde{U}_*\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}}. \quad (43)$$

Then Lipschitzness of the head and Lemma B.7 give

$$\|\tilde{U}_n - \tilde{U}_*\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}} \leq L_V C_{\text{drift}}(n), \quad (44)$$

$$\|\tilde{U}_n - V^{\bar{\pi}}\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}} \leq L_V C_{\text{drift}}(n) + \frac{\bar{\varepsilon}_{\text{res},*}}{1 - \gamma^K}, \quad (45)$$

$$|\bar{\varepsilon}_{\text{res},n} - \bar{\varepsilon}_{\text{res},*}| \leq (1 + \gamma^K) L_V C_{\text{drift}}(n). \quad (46)$$

These are optional comparisons to a memoryless fixed-point reference; none is needed for Corollary B.1.

C DIAGNOSTICS AND PRACTICAL PROXIES

Table 3 shows the theoretical bound quantities and the practical proxies typically logged during training.

Table 2: Claim status. “Measured” denotes retained finite-sample evidence, not a uniform certificate.

Claim	Required condition	Evidence / status
Finite-depth value	Frozen policy pair, common closure and norm	Proven; Theorem 6.1
Persistent direct bound	Complete augmented state, shared frozen map, uniform residual	Valid but uncalibrated; residual not measured uniformly
CPI headline	Exact centering and exact deployed mixture	Not covered; headline policy is distilled
Persistent / PPO results	Nominal implementation protocol	UPI per-seed reevaluation supplied; PPO per-seed run artifact unavailable; interactions unequal
Projection / clamping	Reported factorial with per-seed provenance	+13.0pp composite / +2.2pp inconclusive; per-seed and test artifacts unavailable
Finite-MDP inequalities	Exact finite-state calculation	101 supplied CSV rows have zero numerical violations; generator and stdout supplied

C.1 FINITE-MDP CERTIFICATE

To check that the residual and CPI inequalities hold as stated in a setting where every quantity is computable, we run a controlled finite MDP. This is not a Sudoku benchmark; it checks the certificate mechanism on a 64-state, 4-action discounted MDP with bounded rewards in $[-1, 1]$ and all quantities $(V^\pi, Q^\pi, A^\pi, d_\pi, \eta(\pi))$ solved by matrix inversion. We define a scalar contractive evaluator by

$$z_{t+1}(s) = U_*(s) + L_z(z_t(s) - U_*(s)), \quad U_*(s) = V^\pi(s) + \xi(s), \quad U_n = z_n,$$

with $z_0(s) = 0$, so both the truncation term and the K -step residuals are known in closed form. From U_n we form the centered estimator $\hat{A}_{U_n}$, set $\pi_{\text{cand}}(\cdot | s) = \text{softmax}(\beta \hat{A}_{U_n}(s, \cdot))$, and compute $\eta(\pi_\alpha)$ and $\hat{L}_\pi(\pi_\alpha)$ for $\pi_\alpha = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$ via matrix inversion.

Figure 2 shows two checks. Left: for $\gamma=0.9$, $K=3$, $\sigma_\xi=0.05$, the value error decays with n and worsens as $L_z \rightarrow 1$, staying below the contractive specialization of Theorem 6.1. Right: on an evaluator-based candidate slice ($L_z=0.2$, $n=16$, $\sigma_\xi=0.1$, $\beta=1$), the surrogate gap $\hat{L}_\pi(\pi_\alpha) - \eta(\pi_\alpha)$ is positive and stays below both the decomposition-based CPI penalty and the tighter $\varepsilon_{A, \text{cand}}$ penalty for every tested α (log scale; the bound is conservative). Corollary 6.2 is directly checkable here because $\|U_n - \mathcal{T}_K^\pi U_n\|_\infty$ is computable. The script checks 80 value/decomposition rows (4 L_z values by 20 depths) and 21 α rows on this selected CPI slice, with zero violations; it does not constitute a full Cartesian sweep over $(L_z, n, K, \sigma_\xi, \alpha)$.

For this revision we reran the supplied `finite_mdp_certificate.py`. The submitted `finite_mdp_summary.csv` contains 80 `value_curve` and 21 `cpi_curve` rows, with both certificate flags true in every row; the submitted `finite_mdp_certificate_validation.log` records stdout.

D SECONDARY AND EXPLORATORY DIAGNOSTICS

The following studies probe mechanisms and implementation behavior. They are not theorem validation and are not used as confirmatory hard-suite evidence unless explicitly stated.

Finite-difference L_V diagnostic for the main comparison. To partially operationalize L_V in the same regime used by the main bound, we measure a local finite-difference proxy on the Exp1 refreeze package: frozen B0 states, episodic- z , $n_{\text{eval}}=8$, and seeds 41–50. Table 4 reports the seed-aggregated estimates. The smallest perturbation scale $\epsilon = 10^{-5}$ is noisier, but over the practical range $\epsilon \in \{3 \times 10^{-5}, 10^{-4}, 3 \times 10^{-4}, 10^{-3}\}$ both conditions stay in the same-order 0.085–0.111 band. This remains a local numerical diagnostic rather than a certified global Lipschitz bound, but it reduces the concern that the comparison is dominated by orders-of-magnitude differences in value-head sensitivity.




Figure 2: Finite-MDP certificate diagnostic ($|\mathcal{S}|=64$, $|\mathcal{A}|=4$). Left: value error $\|U_n - V^\pi\|_\infty$ decays with evaluator depth n and stays below the decomposition bound; larger L_z slows the decay. Right: for $\pi_\alpha = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$, the surrogate gap (blue) stays below both the decomposition-based CPI penalty (red) and the tighter $\varepsilon_{A,\text{cand}}$ penalty (green); log scale. All quantities are computed exactly by finite-state dynamic programming.

Table 3: Theory quantities (used in bounds) and practical diagnostics/proxies typically logged during training.

Quantity	Appears in	Practical proxy / diagnostic
L_z^{pre}	Assump. 5.1	Finite-diff Lipschitz of f_θ before projection
L_z^{post}	Assump. 5.1	Finite-diff Lipschitz of $\tilde{f}_\theta = \Pi_R \circ f_\theta$ after projection
$L_z^{\text{sat}}(R)$	Prop. M.3	$(R/\rho_R) L_z^{\text{pre}}(R)$; effective modulus in the saturated regime
ρ_R	Prop. M.3	min-batch $\ f_\theta(z, y, x)\ _2$ for $z \in B_R$; saturation margin
projection_active	App. D.7	Fraction of samples where $\ f_\theta(z)\ > R$ (clipped)
C_z	Assump. 5.1	Empirical $\ \tilde{f}_\theta(z_{\text{init}}, y, x) - z_{\text{init}}\ $ on batches
$L_V(R)$	Eq. (28), Prop. M.4	Lipschitz of V_ψ on B_R ; finite-diff check (Table 4)
ε_0	Prop. M.4	K -step TD error of the boundary-normalized zero-latent baseline B_0
$\ U_n - \mathcal{T}_K^\pi U_n\ $	Proxy only	Held-out K -step TD error; proxy for $\varepsilon_{\text{res}}^*$ up to $O(\ U_n - U_*\)$
$\tilde{\varepsilon}_{\text{res},n}$	Cor. B.1	Finite-batch augmented residual; not measured uniformly for head-line
$\varepsilon_{\text{cent}}$	Thm. 6.5	$\sup_s \mathbb{E}_{a \sim \pi} \hat{A}(s, a) $ (exactly 0 if baseline is exact)
$\varepsilon_{A,\text{cand}}$	Thm. 6.5	Empirical bias of $\mathbb{E}_{a \sim \pi_{\text{cand}}} [\hat{A}]$ on replay states
$\text{KL}(\pi_{\text{new}} \ \pi_\phi)$	App. N	Distillation projection gap measured on replay states

D.1 SPECTRAL NORMALIZATION INSTABILITY AND A MORE ROBUST ALTERNATIVE

The contraction assumption (Assumption 5.1) is a property of the latent update map $z \mapsto \tilde{f}_\theta(z, y, x)$. We first tried to target it by applying power-iteration spectral normalization (Miyato et al., 2018) to linear layers in the inner recursion, followed by a strict scalar factor. On our TRM inner module this turned out to be numerically unstable.

We estimate a local Lipschitz proxy $\hat{L}_z$ by finite differences: sample a small perturbation δ with $\|\delta\|_2 = \varepsilon$ and compute $\|\tilde{f}_\theta(z + \delta, y, x) - \tilde{f}_\theta(z, y, x)\|_2 / \varepsilon$, averaging over random draws. We compare six variants: (i) OFF (no normalization, no scaling), (ii) SN-only, (iii) SCALE-only (only explicit per-layer output scaling), (iv) SN+SCALE, (v) CLAMP-only (operator-norm clamping without scaling), and (vi) CLAMP+SCALE.

Naively applying power-iteration spectral normalization to this architecture does not give reliable contraction control of the inner recursion. Explicit operator-norm clamping stays numerically stable in this check, but the resulting local proxy does not establish Assumption 5.1.

Table 4: Finite-difference $\hat{L}_V$ estimates for the theorem-facing Exp1 refreeze isolation package on the frozen B0 batch at $n_{\text{eval}}=8$. Each entry is mean $\pm$ sample-std over 10 seeds of the per-seed average across 8 repeated random-direction trials (16 perturbations per trial).

Condition	ϵ	$\hat{L}_V$	Seed range
Model A' (clamping disabled)	1e-05	0.192 $\pm$ 0.050	[0.143, 0.274]
Model A' (clamping disabled)	3e-05	0.111 $\pm$ 0.012	[0.095, 0.127]
Model A' (clamping disabled)	1e-04	0.095 $\pm$ 0.009	[0.082, 0.106]
Model A' (clamping disabled)	3e-04	0.096 $\pm$ 0.011	[0.077, 0.110]
Model A' (clamping disabled)	1e-03	0.097 $\pm$ 0.008	[0.080, 0.107]
Model B (clamping enabled)	1e-05	0.200 $\pm$ 0.033	[0.143, 0.250]
Model B (clamping enabled)	3e-05	0.108 $\pm$ 0.023	[0.075, 0.139]
Model B (clamping enabled)	1e-04	0.090 $\pm$ 0.025	[0.056, 0.135]
Model B (clamping enabled)	3e-04	0.087 $\pm$ 0.025	[0.051, 0.122]
Model B (clamping enabled)	1e-03	0.085 $\pm$ 0.026	[0.050, 0.131]

Table 5: Representative local Lipschitz proxy $\hat{L}_z$ for the inner latent update map under different contraction-oriented components (4×4 Sudoku batch; $\varepsilon = 10^{-3}$). Power-iteration spectral normalization (SN) can cause the proxy to explode in our inner recursion, whereas operator-norm clamping (CLAMP) remains stable. These local values are not global certificates.

Variant	$\hat{L}_z$ (proxy)	Status
OFF	≈ 0.86	stable
SN-only	$\approx 2.7 \times 10^4$	exploded
SCALE-only	≈ 0.85	stable
SN+SCALE	$\approx 4.5 \times 10^4$	exploded
CLAMP-only	≈ 0.6	stable
CLAMP+SCALE	≈ 0.6	stable

D.2 LATENT COLLAPSE: ISOLATING $z \rightarrow z$ CLAMPING FROM VALUE-HEAD NORMALIZATION

A natural hypothesis for why $z \rightarrow z$ clamping can hurt success rates is *latent collapse*: the intervention might drive all latent representations toward zero and destroy discriminative information. We test this with a 2×2 diagnostic that isolates $z \rightarrow z$ clamping from value-head normalization, which also affects $\text{Var}(V)$ but is architecturally separate.

The setup is a 2×2 factorial on the 4×4 Sudoku dataset (1–4 empties, single seed, unroll depth $n=16$, batch size 128, projection radius $R=10$):

- **Axis 1 ($z \rightarrow z$ clamping):** disabled vs. enabled (target 0.9).
- **Axis 2 (value-head normalization):** OFF (baseline) vs. ON (target $L_V = 1.0$).

This gives four conditions: A (both OFF), B ($z \rightarrow z$ ON, value-head OFF), C ($z \rightarrow z$ OFF, value-head ON), D (both ON). We measure pre-projection latent norms $\|z_{\text{pre}}^{(n)}\|$ and total variance $\text{Var}(z_{\text{pre}}^{(n)})$, pairwise cosine similarity, value variance $\text{Var}(V)$, and post-projection norms $\|z_{\text{post}}^{(n)}\| = \|\Pi_R(z_{\text{pre}}^{(n)})\|$.

Three observations stand out. First, clamping is associated with modest clustering, not collapse-to-zero. Comparing B vs. A, total variance decreases by $\sim 25\%$ (ratio 0.75) and pairwise cosine similarity increases by ~ 0.21 , while pre-projection norms remain $\|z_{\text{pre}}^{(n)}\| \approx 32$, well above the projection radius $R=10$. Second, the large $\text{Var}(V)$ differences come from value-head normalization, not $z \rightarrow z$ clamping. Comparing C vs. A, $\text{Var}(V)$ increases by $\sim 160,000\times$ (from 0.016 to 2493) when value-head normalization is enabled, while B vs. A barely moves (ratio ≈ 1.04); earlier observations of exploding $\text{Var}(V)$ in a clamping-enabled condition were confounded by value-head normalization. Third, projection saturation explains the constant post-projection norms: all conditions have pre-projection $\|z_{\text{pre}}^{(n)}\| \approx 32 \gg R=10$, so Π_R is always active and clips to $\|z_{\text{post}}^{(n)}\| = 10$.

Table 6: Latent collapse diagnostic at $n=16$ unroll depth. Comparing B vs. A isolates the $z \rightarrow z$ clamping contrast; comparing C vs. A isolates the value-head normalization contrast. Pre-projection metrics are measured before the ball projection Π_R ; post-projection norms are $\|z_{\text{post}}^{(n)}\| = \|\Pi_R(z_{\text{pre}}^{(n)})\|$.

Cond.	$z \rightarrow z$	V-head	$\ z_{\text{pre}}^{(n)}\ $	$\text{Var}(z_{\text{pre}}^{(n)})$	CosSim	$\text{Var}(V)$	$\ z_{\text{post}}^{(n)}\ $
A	OFF	OFF	32.0	814.0	0.196	0.016	10.0
B	ON	OFF	32.0	608.7	0.403	0.016	10.0
C	OFF	ON	32.0	814.0	0.200	2493	10.0
D	ON	ON	32.0	608.7	0.396	2412	10.0

Table 7: Toy-suite feasibility (4×4 Sudoku, 1–4 empties): success rate after 5000 training steps. Mean $\pm$ std over 3 seeds.

Method	Success (%)
UPI-TRM (persistent- z , clamping disabled, $R=10$)	93.3 ± 2.3
UPI-TRM (episodic- z , clamping disabled, $R=10$)	93.3 ± 2.3
UPI-TRM (episodic- z , clamping enabled, proj. off)	90.7 ± 5.0
Random (masked uniform)	52.0

Within this single-seed diagnostic, the modest clustering pattern is inconsistent with a simple collapse-to-zero explanation, but it cannot identify the mechanism behind success-rate differences.

A 200-step follow-up associates value-head normalization with clamp-hitting: clamping alone remains stable with targets ≈ -1.7 , value-head normalization alone reaches the ± 20 clamps, and the combined cell reaches the positive clamp (target = 20 by step 60, $V(s) = 20$ by step 100). Short clamping-only runs with value-head normalization disabled remain stable but obtain only $\approx 10\%$ success at 1k steps. These exploratory observations identify a confound; they do not establish a causal or general training mechanism.

D.3 TOY-SUITE SANITY CHECK


The 1–4-empty toy suite is a sanity check because masked random already succeeds 52% of the time; the depth-mismatch results reported on it (Appendix F.1) are calibration only. The full success-rate table is recorded here for completeness.

D.4 ISOLATION STUDY: DETAILED TABLES


Tables 8 and 9 provide the full numeric results for the B0 isolation study (Section F.1).

D.5 FINITE- R SWEEP ON B0 (FIXED PAIR)


To test whether the finite- R specialization in Remark M.2 has empirical contact with the evaluator-truncation channel, we reran the fixed B0 pair $(n_{\text{train}}, n_{\text{eval}}) = (2, 8)$ from Section F.1 over radii $\{\text{proj. off}, 5, 10, 20, 30, 32, 34, 40, 100\}$ and aggregated seed-level means over seeds 41–50. The fixed-pair analysis passes the endpoint-separation and inactive-plateau checks, the plateau is recovered exactly by $R \geq 40$, and the source pre-projection norm at $R=10$ lies in a narrow p_{05} – p_{95} band of about 30.6–33.2, which is why the informative transition window is narrow. Across both models, $R=20$ is the cleanest shared intermediate regime. Adding held-out theory summaries shows that the measurable Bellman residual moves in that same narrow window, but non-monotonically in R : it drops from the inactive plateau to a low point around $R=10$, rises again through the intermediate window, and returns to the inactive plateau for $R \geq 40$. This gives finite- R contact with the residual channel as well as the truncation channel, while still not identifying $\varepsilon_{\text{res}}^*(R)$ end-to-end.



(a) Primary stability metrics.



(b) Mechanism diagnostics.



(c) Held-out theory diagnostics.

Figure 3: Fixed-pair finite- R sweep on B0 initial states with $n_{\text{train}}=2$, $n_{\text{eval}}=8$, and seeds 41–50. Top: the primary metrics (Δ_V , Δ_z , argmax agreement, Δ_π) move smoothly between the active-projection regime and the inactive plateau, with $R=20$ as the cleanest shared intermediate exhibit and $R \geq 40$ matching *proj. off*. Middle: the mechanism diagnostics show why this is not a gradual saturation knob. The saturated rate stays at 1.0 through $R \leq 32$, drops to 0.2 at $R=34$, and to 0.0 at $R \geq 40$, while *proj_disp_proxy* is only a scalar proxy based on norms, not a displacement norm. Bottom: held-out theory diagnostics show that the measurable Bellman residual also moves in the same transition window, but non-monotonically in R , while $\hat{L}_V$ remains same-order across radii. This appendix result therefore provides bounded empirical contact for both the truncation channel and the measurable residual channel, but it still does not identify any R -dependence of $\varepsilon_{\text{res}}^*$ end-to-end or explain projection dominance on the hard suite.

Table 8: Unroll sensitivity on B0 (initial states). Training depth $n_{\text{train}}=2$, evaluated at fixed $4\times$ depth ($n_{\text{eval}}=8$).

Condition	Δ_V	Δ_π	Δ_z	Argmax
Clamping disabled	0.166 ± 0.287	0.0002 ± 0.0006	4.26 ± 1.53	0.902 [0.88,0.92]
Clamping enabled	0.064 ± 0.060	0.0001 ± 0.0002	1.67 ± 0.28	0.955 [0.94,0.97]

Table 9: Projection radius sweep on B0 (initial states), using the fixed $n = 2 \rightarrow 8$ mismatch. Here *proj. off* denotes $\Pi_R = \text{Id}$; the clamping-enabled condition has a $2.1\times$ lower value-instability proxy.

Radius	Condition	Δ_V	Δ_z	Argmax
proj. off	Clamping disabled	2.275 ± 3.413	36.1 ± 7.9	0.47 [0.44,0.50]
	Clamping enabled	1.097 ± 0.912	22.6 ± 3.3	0.65 [0.62,0.68]
R=10	Clamping disabled	0.166 ± 0.287	4.3 ± 1.5	0.90 [0.88,0.92]
	Clamping enabled	0.064 ± 0.060	1.7 ± 0.3	0.95 [0.94,0.97]
R=100	Clamping disabled	2.275 ± 3.413	36.1 ± 7.9	0.47 [0.44,0.50]
	Clamping enabled	1.097 ± 0.912	22.6 ± 3.3	0.65 [0.62,0.68]

Note: The $R = 100$ rows match the *proj. off* rows because pre-projection latent norms stay well below 100, so clipping never activates at this radius.

D.6 ISOLATION STUDY: B1 SUCCESSOR STATES

We present the full B1 (one-step successor states from B0) results from the isolation study (Section F.1).

Unroll sensitivity on B1. On successor states, the clamping-enabled condition has lower Δ_V , Δ_z , and Δ_π and higher action agreement at all tested depth mismatches, though absolute instability remains higher than on B0 initial states.

Radius sweep on B1. On B1 at *proj. off*, the clamping-enabled condition has lower latent drift (Δ_z : $35.3 \rightarrow 22.7$, i.e., $1.6\times$), higher action agreement ($47.2\% \rightarrow 64.7\%$), and lower value instability Δ_V ($2.23 \rightarrow 1.06$, i.e., $2.1\times$). Absolute variability nevertheless remains substantially higher than on B0 initial states.

The main-text claim in Section F.1 stays anchored on B0 initial states, because B1 is an appendix-only successor-state check and is noisier even though the qualitative direction now matches B0.

D.7 PROJECTION DOMINANCE STUDY

We probed the clamping target $L_z^{\text{target}} \in \{0.9, 0.95, 0.99, 0.999\}$. All models are trained on the same 4×4 Sudoku suite (1–4 empties) with `disable_value_head_norm: true`, evaluated at $8\times$ depth mismatch on B0 (initial states).

The target does not control the measured local proxy. Table 11 shows that $\hat{L}_z^{\text{post}}$ stays in a narrow 0.219–0.242 band across the four target settings, and the ordering does not produce a monotonic ordering in stability metrics. These finite-difference values are not certified global Lipschitz constants.

Projection, separately, is a major stabilizer in this exploratory comparison. Table 12 shows that projection at $R=10$ improves Δ_V by 6–10 $\times$ and argmax agreement by ≈ 8 –10 pp compared to *proj. off*. At $R=10$ projection is active 100% of the time, which can mask differences between the clamping conditions.

In this projection-active regime, the clamping target does not meaningfully vary the local proxy, while projection itself acts as a strong stabilizer. The isolation study therefore supports only a binary clamping-condition contrast at *proj. off*, not a tunable or certified post-projection modulus.

Table 10: Radius sweep on B1 (successor states). Under the refrozen 10-seed protocol, the clamping-enabled condition has better values for all three metrics with *proj. off*, though absolute instability remains higher than on B0.

Radius	Condition	Δ_V	Δ_z	Argmax
proj. off	Clamping disabled	2.230±3.477	35.3±8.7	0.47 [0.46,0.48]
	Clamping enabled	1.064±0.929	22.7±3.5	0.65 [0.64,0.66]
R=10	Clamping disabled	0.218±0.396	4.6±1.8	0.89 [0.89,0.90]
	Clamping enabled	0.066±0.060	1.7±0.3	0.96 [0.95,0.96]
R=100	Clamping disabled	2.230±3.477	35.3±8.7	0.47 [0.46,0.48]
	Clamping enabled	1.064±0.929	22.7±3.5	0.65 [0.64,0.66]

Note: The $R = 100$ rows match the *proj. off* rows because pre-projection latent norms stay well below 100, so clipping never activates at this radius.

Table 11: Spectral-norm dial does not control Lipschitz constant. At $R=10$, projection dominates (100% active) and $\hat{L}_{post}$ saturates. With *proj. off* ($\Pi_R = \text{Id}$), $\hat{L}_{pre}$ varies weakly (spread $0.064 < 0.08$ threshold) and non-monotonically.

Target L_z	$\hat{L}_{pre}$ (<i>proj. off</i>)	$\hat{L}_{post}$ ($R=10$)	Proj Active ($R=10$)	Proj Active (<i>proj. off</i>)
0.9	0.473±0.002	0.242±0.003	100%	0%
0.95	0.409±0.029	0.219±0.005	100%	0%
0.99	0.448±0.016	0.234±0.003	100%	0%
0.999	0.440±0.044	0.231±0.010	100%	0%
Spread	0.064	–	–	–

D.8 PROJECTION-FREE DIAL RANGE TEST

Appendix D.7 shows that the spectral-norm target L_z^{target} is not an effective control parameter when projection is active. We therefore test a narrower question here: with projection disabled, can inference-time scaling induce a wider range in the local proxy?

With projection disabled (*proj. off*, implemented as $\Pi_R = \text{Id}$) and value-head spectral norm disabled, we apply inference-time scalar scaling to the $z \rightarrow z$ update weights of a single trained checkpoint, then measure the achieved pre-projection Lipschitz proxy L_z^{pre} and stability metrics (Δ_V , argmax agreement) at $8\times$ depth mismatch on 100 B0 samples.

The achieved L_z^{pre} span is wide ($\approx 0.52\text{--}0.95$; spread ≈ 0.43), so inference-time scaling *can* vary the proxy once projection is disabled. What is not supported is a *monotonic* stability relation: neither Δ_V nor argmax agreement changes monotonically with L_z^{pre} in this range. We therefore treat the run as a feasibility check only.

D.9 TRAINING STABILITY ABLATION

We ablate projection at *training time* to test whether the ball projection Π_R is necessary for stable training, as opposed to its role in evaluation-time mismatch stability (see projection dominance study above).

In this easy-suite ablation, all six conditions (clamping enabled/disabled crossed with {*proj. off*, 10, 100}) pass the stated stability gate and obtain 94–98% success. Thus projection is not necessary for stability within this tested protocol; this is not a general training-stability conclusion.

D.10 PROJECTION-FREE CLAMPING-TARGET SWEEP (FULL RESULTS)

The depth-mismatch summary in Appendix F.1 gives the high-level takeaway. Table 15 below provides the full numeric results at $8\times$ depth mismatch.

Table 12: Projection is the primary stabilizer. $R=10$ projection dramatically improves stability metrics compared to *proj. off* ($\Pi_R = \text{Id}$).

Target L_z	ΔV (n=2→16)		Argmax Agree (%)	
	$R=10$	<i>proj. off</i>	$R=10$	<i>proj. off</i>
0.9	2.18	10.48	99.0	90.7
0.95	0.55	9.95	97.3	89.3
0.99	0.99	7.72	99.7	88.0
0.999	1.58	10.02	99.3	89.7
Mean	1.32	9.54	98.8	89.4
Range	[0.5, 2.2]	[7.7, 10.5]	[97, 100]	[88, 91]

Table 13: Projection-free dial range test via inference-time scaling. Single checkpoint with *proj. off*, value-head spectral normalization disabled, 100 eval samples at $8\times$ depth mismatch. The achieved L_z^{pre} proxy spans $\approx 0.52\text{--}0.95$.

Scale	L_z^{pre}	$\Delta_V @ 8\times$	Argmax @ $8\times$
1.00	0.9483	0.5663	0.6600
0.90	0.7192	0.4918	0.6600
0.80	0.5876	0.4174	0.7000
0.70	0.5195	0.4389	0.7400
0.60	0.5212	0.6385	0.6900

D.11 STABILITY-EXPRESSIVITY TRADEOFF (FULL RESULTS)

This appendix provides an exploratory stability-expressivity analysis. **Important:** Success rates here ($\approx 6.7\%$) are *not* comparable to Table 7 ($\approx 93\%$) because they use different evaluation protocols: Exp5 uses *matched compute* ($n_{\text{eval}}=2$, *proj. off*) to isolate clamping-condition differences, while Table 7 uses $n_{\text{eval}}=2$ with mixed projection settings (clamping disabled: $R=10$; clamping enabled: *proj. off*). The low matched-compute success in Exp5 reflects the inherent difficulty of solving puzzles with only 2 unroll steps under uniform projection-off evaluation, not model deficiency.

Using inference-time scaling (scales $\in \{1.0, 0.85, 0.70, 0.55\}$, *proj. off*, value-head spectral normalization disabled), we observe that stability varies modestly across scale settings while task success remains essentially flat ($\approx 6.7\%$); we do not observe a pronounced tradeoff in this regime. Table 16 provides the numeric results.

E DECISION-RELEVANT HARD-SUITE RECORDS

This section retains the hard-suite task records and reported projection $\times$ clamping aggregates used in the main narrative.

E.1 HARDER 4×4 SUDOKU (6–8 EMPTIES)

We extend evaluation to a harder suite with 6–8 empty cells, where the random-policy success rate drops to 0%.

E.2 REPORTED PROJECTION $\times$ CLAMPING AGGREGATES ON HARD 4×4

The manuscript records aggregate cells from a 2×2 factorial on the same no-mask hard 4×4 suite. The per-seed factorial artifact and statistical test script are unavailable, so the table supports only arithmetic comparisons among the reported aggregates.

The reported projection contrast raises success from 36.2% to 49.2% (+13.0pp averaged over clamping status). It is a composite algorithmic effect spanning the forward pass and training dynamics

Table 14: Training-time projection ablation (4×4 trivial Sudoku, feasibility checker). Projection setting $\in \{\text{proj. off}, 10, 100\}$ crossed with clamping enabled/disabled. All conditions pass the stated stability gate. Within this tested easy-suite protocol, success rates are comparable with *proj. off*; this is not a general training-stability claim.

Condition	Clamping	R	Success (%)	Stability Gate
No-Contr / $R=10$	OFF	10	94	PASS
No-Contr / $R=100$	OFF	100	96	PASS
No-Contr / <i>proj. off</i>	OFF	<i>proj. off</i>	96	PASS
Contr / $R=10$	ON	10	98	PASS
Contr / $R=100$	ON	100	96	PASS
Contr / <i>proj. off</i>	ON	<i>proj. off</i>	96	PASS

Table 15: Projection-free clamping-target results with *proj. off* and value-head spectral normalization disabled. Monotonic correlation with scale setting is statistically significant (Spearman ρ shown below); the measured $\hat{L}_z$ remains a local proxy.

Scale	$\hat{L}_{\text{preproj}}$	Argmax ($n_{\text{eval}}=8$) B0	Argmax ($n_{\text{eval}}=8$) B1	ΔV ($n_{\text{eval}}=8$) B0	Entropy
1.00	0.470 ± 0.326	0.887 ± 0.133	0.856 ± 0.156	0.267	2.61
0.85	0.383 ± 0.173	0.867 ± 0.148	0.862 ± 0.129	0.176	2.60
0.70	0.391 ± 0.084	0.880 ± 0.118	0.860 ± 0.102	0.532	2.65
0.55	0.490 ± 0.037	0.820 ± 0.080	0.793 ± 0.091	1.336	2.81
Spearman ρ (B0 Argmax, $n_{\text{eval}}=8$): $\rho = -0.866$ [95% CI: -1.000, -0.542]					
Spearman ρ (B1 Argmax, $n_{\text{eval}}=8$): $\rho = -0.923$ [95% CI: -1.000, -0.787]					

(changes to f_θ , ε_0 , and the optimizer trajectory), which (70) does not model. Averaging over projection status, the clamping-enabled contrast raises success from 41.6% to 43.8% (+2.2pp). The source-reported $p \approx 0.6$ is not independently auditable because the retained package lacks the per-seed factorial artifact and test script; the aggregate contrast is inconclusive and does not support a variance-reduction claim.

F SECONDARY ROBUSTNESS DIAGNOSTICS

These depth-mismatch curves and discussion details are secondary diagnostics and do not establish hard-suite task performance.

F.1 DEPTH-MISMATCH ROBUSTNESS (ISOLATION STUDY)

To probe evaluation-time robustness as compute changes, we evaluate checkpoints trained at $n_{\text{train}}=2$ under deeper unrolls $n_{\text{eval}} \in \{4, 8, 16\}$ on the 1–4-empty suite and aggregate over seeds 41–50. On B0 initial states at fixed $4 \times$ mismatch ($n=2 \rightarrow 8$), the clamping-enabled condition lowers Δ_V from 0.166 to 0.064 and raises argmax agreement from 90.2% to 95.5%; Δ_π stays near zero. This is finite-batch diagnostic contact with the episodic setting, not theorem validation, and it does not establish the global contraction condition or speak to persistent- z behavior on the hard suite.

The detailed tables (Appendix D.4) show that moving from *proj. off* to $R=10$ yields the larger reduction in Δ_V , while Appendix D.6 gives the same qualitative pattern on B1 successor states with higher noise.

F.2 REPORTED 2×2 (CLAMPING $\times$ PROJECTION) ON HARD SUITE

Historical hard-suite comparisons mixed projection settings and therefore confounded the clamping intervention with clipping. The recorded 2×2 aggregate uses the same no-mask hard 4×4 Sudoku protocol (6–8 empties, 20k steps, 10 seeds), but the submitted evidence lacks per-seed provenance.

Three readings follow from the aggregate table. Enabling projection raises reported mean success from 36.2% to 49.2% (+13.0pp averaged over clamping status), a composite algorithmic contrast. Clamping adds only a small average contrast: at *proj. off* it changes success from 35.0% to 37.4%

Table 16: Exp5: Stability–Expressivity Analysis. Success remains essentially flat across scale settings ($\approx 6.7\%$); no clear tradeoff observed in this regime.

Scale	$\hat{L}_{\text{preproj}}$	Stability (argmax@8 $\times$)	Success (trivial)	$\Delta V@8\times$
1.00	0.801 ± 0.238	0.700 ± 0.085	0.067 ± 0.017	0.528
0.85	0.562 ± 0.136	0.730 ± 0.057	0.070 ± 0.014	0.439
0.70	0.489 ± 0.073	0.770 ± 0.098	0.067 ± 0.017	0.458
0.55	0.524 ± 0.031	0.730 ± 0.107	0.067 ± 0.012	0.680

Table 17: Manuscript-reported no-mask hard 4×4 2×2 aggregates (10 seeds per cell). The projection contrast is larger; the clamping-enabled contrast is small and inconclusive. Per-seed artifacts are unavailable.

Condition	Success (%)
Clamping disabled, proj. off	35.0 ± 8.9
Clamping disabled, $R=10$	48.2 ± 7.1
Clamping enabled, proj. off	37.4 ± 7.8
Clamping enabled, $R=10$	50.2 ± 10.5

(+2.4pp), and at $R=10$ from 48.2% to 50.2% (+2.0pp). The highest-mean cell is clamping enabled with $R=10$ ($50.2\% \pm 10.5\%$) versus $48.2\% \pm 7.1\%$ for clamping disabled at $R=10$; these aggregates do not establish a contraction benefit or variance reduction.

The reported table replaces the older asymmetric narrative: projection has the larger composite success contrast, while clamping has only a small, inconclusive average contrast. The separate depth-mismatch result concerns value stability, not hard-suite return.

F.3 DISCUSSION DETAILS

The stability experiments answer different questions, and the shortest consistent reading is: Appendix F.1 gives a binary clamping-condition contrast on *evaluation-time* depth-mismatch stability for B0 initial states; Appendix D.7 shows that active projection can dominate and mask that effect; Appendix D.9 shows projection is not required for *training* stability; and Appendix D.10 shows that once projection is disabled, the local proxy regains meaningful range. Appendix D.11 is exploratory only: its success rates ($\approx 6.7\%$) use matched compute ($n=2$, proj. off) and are not directly comparable to Table 7, which mixes projection settings.

These appendices show two observed contrasts: projection changes success and depth-stability diagnostics through coupled forward-pass and training channels, while disabling projection exposes a wider local-proxy range for inference-time checks. They do not isolate a projection training-stability effect, certify contraction, or establish that clamping improves hard-suite performance.

A naive reading of the older asymmetric comparison would suggest that the clamping intervention is strongly harmful on harder tasks. The reported no-mask 10-seed 2×2 aggregates do not support that reading. Projection changes mean success from 36.2% to 49.2%, while clamping changes it only from 41.6% to 43.8% averaged over projection settings (+2.2pp) and does not reduce variance under projection. Mathematical contraction is a sufficient error-control condition, but these data do not establish that condition or demonstrate clamping as a hard-task performance dial.

G PROTOCOL AND THEORY-ALIGNED RUN RECORDS

This section records the interaction-budget audit, headline significance data, and the theory-aligned episodic hard-suite run. The common-budget baseline table lacks standalone provenance in the retained package and is marked accordingly.

Table 18: Artifact status for the missing interaction-equalized comparison.

Run family	UPI-TRM arm	Standalone prove- nance	Claim status
Source baseline-only audit	Absent	Absent	Numerical rows ex- cluded
Required UPI-TRM vs. PPO audit	Required	Not run	Missing experiment

G.1 INTERACTION-EQUALIZED BASELINE AUDIT

Table 1 reports the nominal 20k outer-step protocol. Because PPO, A2C, and DQN collect 64, 32, and 4 transitions per outer step, it does not equalize interactions. The source manuscript recorded baseline-only common-budget rows, but their standalone logs are absent. We therefore remove those unauditable numbers from active claims (they remain recoverable in Git history) and retain only the artifact-status audit below.

There is therefore no numerical interaction-matched UPI-TRM versus PPO result in this paper.

G.2 HEADLINE SIGNIFICANCE-TEST DATA

The significance calculation in Section 7.1 uses the ten seed-level success rates recorded in the manuscript for the nominal 20k outer-step hard-suite protocol. For UPI-TRM, the rates are also present in the submitted reevaluation artifact:

0.56, 0.50, 0.68, 0.50, 0.64, 0.70, 0.74, 0.42, 0.38, 0.62,

with mean 57.4% and sample standard deviation 12.2%. For the architecture-matched TRM+PPO baseline, the corresponding rates are

0.30, 0.32, 0.12, 0.54, 0.12, 0.56, 0.20, 0.32, 0.42, 0.30,

with mean 32.0% and sample standard deviation 15.3%. Applying Welch’s t -test to these recorded values gives $t \approx 4.10$, $df \approx 17.1$, two-sided $p \approx 7.3 \times 10^{-4}$, and a 95% confidence interval of [12.4, 38.5] percentage points for the mean difference. These calculations establish arithmetic consistency only: the submitted evidence does not contain the PPO per-seed run artifact or an independent resampling script.

G.3 THEORY-ALIGNED EPISODIC- z RERUN ON THE HARD SUITE

We ran a protocol-specified, theory-aligned rerun on the same no-mask hard 4×4 Sudoku suite used in Table 1 (6–8 empties, $T=16$, 20k environment steps, seeds 41–50). The supplied files do not establish that the protocol was timestamped before execution. The rerun starts from the persistent- z hard-suite configuration and jointly changes four factors: episodic- z latents, exact baseline summation, exact K -step targets, and explicit conservative mixtures. The architecture-matched evaluation interface and the inherited persistent- z hyperparameters ($\gamma=0.99$, $K=1$, $n=2$) are unchanged. Diagnostics are computed on a frozen 256-instance closure batch (seed 1729). The fixed test split of the canonical 6–8-empty dataset contains only 99 viable puzzles, so the closure batch is sampled from the `train/` split with an independent RNG path; this deviation affects only the diagnostic batch, not the task-evaluation pool or the training seeds.

The paper tables below curate the main quantities. The submitted supplement includes `EPISODIC_Z_RUN_LOG.md`, `episodic_diagnostics_summary.json`, `cpi_penalty_grid.json`, and `episodic_per_seed_diagnostics.jsonl`.

Table 19 reports the final task outcome and the main diagnostics at 20k. The rerun reaches 0.0% success over 10 seeds, so by the locked decision rule it falls into Branch D: $S_E = 0.0\% < 0.85 \times 57.4\% = 48.79\%$. The persistent- z value $S_P = 57.4\%$ in Table 1 is measured under the nominal 20k outer-step protocol of Section 7.1, which uses a larger environment-interaction budget than this rerun; the Branch D margin holds under either budget framing. We therefore retain persistent- z as the headline empirical result and use the episodic- z rerun only for finite-batch diagnostics on the frozen closure batch.

Table 19: Protocol-specified theory-aligned episodic- z rerun at 20k on the no-mask hard 4×4 suite. Means $\pm$ sample std over 10 seeds unless noted.

Quantity	Value
Success (%)	0.0 ± 0.0
Return	-21.733 ± 0.512
Score	3.906 ± 0.606
Invalid	0.000 ± 0.000
Decision branch	D ($0.0\% < 0.85 \times 57.4\% = 48.79\%$)
$\varepsilon_{\text{cent}}$ @ 20k	$(1.782 \pm 0.301) \times 10^{-6}$
$\varepsilon_{\text{res},n}$ @ 20k	5.047 ± 0.175
projection-active rate @ 20k	1.0
ρ_R @ 20k	$45.2547 \pm 1.3 \times 10^{-5}$
finite-seed $\hat{L}_z^{\text{post}}$ @ 20k	0.600 ± 0.104 over 8/10 seeds
saturating seeds	41, 46 ($\hat{L}_z^{\text{post}} \geq 1$)

Table 20: Predicted CPI penalty at 20k on the episodic- z closure batch. Finite-subset mean $\pm$ sample std is computed over seeds 42, 43, 44, 45, 47, 48, 49, and 50. The all-seed mean is $+\infty$ at every α because seeds 41 and 46 have $\hat{L}_z^{\text{post}} \geq 1$.

α	finite-subset predicted penalty	all-seed mean
0.05	5215.11 ± 164.00	$+\infty$
0.10	10809.78 ± 303.85	$+\infty$
0.20	23137.80 ± 561.71	$+\infty$
0.40	52348.56 ± 1387.43	$+\infty$

At 20k, the finite-batch centering defect is numerically small, projection is active on every sampled closure-batch state, and the held-out finite-depth Bellman residual remains large. None of these finite-batch observations is a uniform certificate. Two seeds (41 and 46) have $\hat{L}_z^{\text{post}} \geq 1$, so the all-seed mean predicted CPI penalty is literally $+\infty$ at every $\alpha \in \{0.05, 0.10, 0.20, 0.40\}$. For interpretability, Table 20 also reports finite-seed summaries over the remaining 8/10 seeds.

H NOTATION SUMMARY

We summarize notation here for convenience. Table 21 lists additional symbols used in proofs.

I CONSTRUCTIVE DEFINITION OF THE ADVANTAGE-EVALUATION CLOSURE

Definition I.1 (Reachable set). For a policy π in a finite or countable MDP, let $\mathcal{R}_\pi$ be the exact set of states reached with positive probability under some state marginal ρP_π^t , $t \geq 0$, where P_π is the policy-induced transition kernel.

Definition I.2 (Advantage-evaluation closure). Fix a current policy π and a candidate policy π_{cand} . Let $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ contain $\mathcal{R}_\pi$ and s_{abs} such that

- for every $s \in \mathcal{R}_\pi$ and every action in $\text{supp}(\pi(\cdot|s)) \cup \text{supp}(\pi_{\text{cand}}(\cdot|s))$, the next state lies in $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ almost surely; and
- for every $s \in \mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ and $a \in \text{supp}(\pi(\cdot|s))$, $P(\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} | s, a) = 1$.

We write $\|f\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}$ for the ordinary supremum of $|f|$ on this set, including the absorbing boundary.

The following construction gives the exact closure in Definition I.2.

Fix a current policy π and a candidate policy π_{cand} . Define the *one-step deviation* set

$$\mathcal{R}_{\pi, \pi_{\text{cand}}}^+ := \mathcal{R}_\pi \cup \{s' : \exists s \in \mathcal{R}_\pi, \exists a \in \text{supp}(\pi(\cdot|s)) \cup \text{supp}(\pi_{\text{cand}}(\cdot|s)) \text{ with } P(s' | s, a) > 0\}. \quad (47)$$

Table 21: Additional symbols and notations

Symbol	Meaning
$z^{(j)}(s), U_j(s)$	Latent and scalar value after finite depth j
$U_*(s)$	Fixed-point value used only in the contractive specialization
$B_{n,m}$	Finite path/discrepancy plus depth- m residual certificate
$\tilde{U}_n(\bar{s}), \tilde{U}_*(\bar{s})$	Persistent finite-depth value / memoryless reference
$\varepsilon_{\text{res},n}$	Episodic finite-depth Bellman residual
$\bar{\varepsilon}_{\text{res},n}$	Persistent augmented-state finite-depth residual
$\bar{\varepsilon}_{\text{res},*}$	Augmented residual of the fixed-point reference
$L_{z^*}, \Delta y_{\text{max}}, E_0$	Slow-drift parameters
$\varepsilon_A, \bar{\varepsilon}_A$	Episodic / augmented advantage error bounds
$\varepsilon_{\text{res}}^*$	Architectural residual at U_*
ε_{CPI}	CPI constant
$\kappa_n = L_z^n$	Contraction after n steps
$\mathcal{R}_{\Pi}$	Union of reachable sets for policy family Π
$\ f\ _{\infty, \Pi}$	Sup-norm over $\mathcal{R}_{\Pi}$
$B_V(s)$	Baseline: $\mathbb{E}_{b \sim \pi}[\hat{Q}_V(s, b)]$

Let $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ denote the smallest superset of $\mathcal{R}_{\pi, \pi_{\text{cand}}}^+$ that is closed under forward dynamics of π :

$$\begin{aligned} \mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} &:= \bigcup_{t \geq 0} \mathcal{R}^{(t)}, \quad \mathcal{R}^{(0)} = \mathcal{R}_{\pi, \pi_{\text{cand}}}^+, \\ \mathcal{R}^{(t+1)} &= \mathcal{R}^{(t)} \cup \{s' : \exists s \in \mathcal{R}^{(t)}, \\ &\quad a \in \text{supp}(\pi(\cdot | s)), P(s' | s, a) > 0\}. \end{aligned} \quad (48)$$

The sup-norm $\|f\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}$ used throughout is the one introduced in Definition I.2.

Remark I.3 (Requirements for a general-kernel extension). A general measurable-state-space extension is not part of the formal claims above. One route is to fix a single sigma-finite reference measure ν on the closure, with $\nu(\{s_{\text{abs}}\}) > 0$, such that every relevant one-step transition law and every K -step π transition law is absolutely continuous with respect to ν . All essential suprema must then be taken in the same $L^\infty(\nu)$ space so that $\mathcal{T}_K^\pi$ is well defined on equivalence classes. For general action kernels, support conditions must be replaced by statements holding for $\pi(\cdot | s)$ - and $\pi_{\text{cand}}(\cdot | s)$ -almost every action unless additional topology is specified. The CPI occupancy argument must likewise be written for finite signed measures in total variation rather than for probability vectors in ℓ_1 .

J SHAPING CONVENTION AND ABSORBING-STATE NORMALIZATION

We use potential-based shaping with potential $\Phi(s) = c(x, y)$ for $s = (x, y)$ and $\Phi(s_{\text{abs}}) = C_{\text{max}}$, and define

$$r(s, a, s') = r_0(s, a, s') + \gamma \Phi(s') - \Phi(s), \quad (49)$$

applying shaping on *all* transitions, including the absorbing self-loop.

For the absorbing self-loop, taking $r_0(s_{\text{abs}}, a, s_{\text{abs}}) = 0$ gives

$$r(s_{\text{abs}}, a, s_{\text{abs}}) = \gamma \Phi(s_{\text{abs}}) - \Phi(s_{\text{abs}}) = (\gamma - 1)C_{\text{max}}.$$

Thus every policy satisfies the fixed shaped boundary value

$$V^\pi(s_{\text{abs}}) = \sum_{t=0}^{\infty} \gamma^t (\gamma - 1) C_{\text{max}} = -C_{\text{max}}.$$

If one prefers the standard normalization $V(s_{\text{abs}}) = 0$, use the shifted potential $\tilde{\Phi}(s) := \Phi(s) - C_{\text{max}}$. This adds the constant reward $(1 - \gamma)C_{\text{max}}$ to every transition and shifts all returns by $+C_{\text{max}}$ independently of the policy. Accordingly, performance differences and optimal policies are unchanged.

K CPI DERIVATION FOR MIXTURE POLICIES

We derive the CPI bound (29) specialized to mixture policies $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$, following the framework of Kakade & Langford (2002).

This derivation is stated for finite or countable state and action spaces, where discounted occupancies are probability mass functions and $\|\cdot\|_1$ has its usual meaning. A general-kernel extension requires the total-variation formulation and common-domination conditions in Remark I.3.

Setup. Define the *advantage function under the greedy policy* π_{cand} :

$$\delta(s) := \mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)}[A^\pi(s, a)]. \quad (50)$$

Note that $|\delta(s)| \leq \varepsilon_{\text{CPI}} := \sup_s |\delta(s)|$ by definition.

For the mixture policy $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$:

$$\mathbb{E}_{a \sim \pi_{\text{new}}(\cdot|s)}[A^\pi(s, a)] = (1 - \alpha) \underbrace{\mathbb{E}_{a \sim \pi}[A^\pi(s, a)]}_{=0} + \alpha \mathbb{E}_{a \sim \pi_{\text{cand}}}[A^\pi(s, a)] \quad (51)$$

$$= \alpha \delta(s). \quad (52)$$

Action-set convention. If the available edit actions depend on the plan (and hence on the state $s = (x, y)$), we interpret $\pi(\cdot|s)$ and $\pi_{\text{cand}}(\cdot|s)$ as distributions on the same per-state action set $\mathcal{A}(s)$. Equivalently, we may embed both policies into a fixed superset action space by assigning probability 0 to invalid actions. Under either convention, $\|\pi'(\cdot|s) - \pi(\cdot|s)\|_1 \leq 2$ holds for all s . **Performance difference lemma.** The standard performance difference lemma (Kakade & Langford, 2002) gives:

$$\eta(\pi_{\text{new}}) - \eta(\pi) = \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\mathbb{E}_{a \sim \pi_{\text{new}}}[A^\pi(s, a)]] = \frac{\alpha}{1 - \gamma} \mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)]. \quad (53)$$

Distribution shift bound. We need to relate $\mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)]$ to $\mathbb{E}_{s \sim d_\pi}[\delta(s)]$. The key is to bound $\|d_{\pi_{\text{new}}} - d_\pi\|_1$.

Lemma K.1 (Occupancy measure coupling). *For any two policies π' and π in an MDP with discount γ :*

$$\|d_{\pi'} - d_\pi\|_1 \leq \frac{\gamma}{1 - \gamma} \mathbb{E}_{s \sim d_{\pi'}}[\|\pi'(\cdot|s) - \pi(\cdot|s)\|_1]. \quad (54)$$

Proof. Recall the discounted occupancy satisfies the fixed-point equation

$$d_\pi = (1 - \gamma)\rho + \gamma d_\pi P_\pi, \quad d_{\pi'} = (1 - \gamma)\rho + \gamma d_{\pi'} P_{\pi'}.$$

Subtracting gives

$$d_{\pi'} - d_\pi = \gamma(d_{\pi'} P_{\pi'} - d_\pi P_\pi) = \gamma(d_{\pi'}(P_{\pi'} - P_\pi) + (d_{\pi'} - d_\pi)P_\pi).$$

Taking ℓ_1 norms and using that P_π is a Markov operator (so it is non-expansive in ℓ_1),

$$\|d_{\pi'} - d_\pi\|_1 \leq \gamma\|d_{\pi'}(P_{\pi'} - P_\pi)\|_1 + \gamma\|(d_{\pi'} - d_\pi)P_\pi\|_1 \leq \gamma\|d_{\pi'}(P_{\pi'} - P_\pi)\|_1 + \gamma\|d_{\pi'} - d_\pi\|_1.$$

Rearranging yields

$$\|d_{\pi'} - d_\pi\|_1 \leq \frac{\gamma}{1 - \gamma} \|d_{\pi'}(P_{\pi'} - P_\pi)\|_1.$$

Finally,

$$\begin{aligned} \|d_{\pi'}(P_{\pi'} - P_\pi)\|_1 &= \left\| \sum_s d_{\pi'}(s) (P_{\pi'}(\cdot|s) - P_\pi(\cdot|s)) \right\|_1 \\ &\leq \sum_s d_{\pi'}(s) \|P_{\pi'}(\cdot|s) - P_\pi(\cdot|s)\|_1 \\ &= \mathbb{E}_{s \sim d_{\pi'}}[\|P_{\pi'}(\cdot|s) - P_\pi(\cdot|s)\|_1] \\ &= \mathbb{E}_{s \sim d_{\pi'}} \left[\left\| \sum_a (\pi'(a|s) - \pi(a|s)) P(\cdot|s, a) \right\|_1 \right] \\ &\leq \mathbb{E}_{s \sim d_{\pi'}} \left[\sum_a |\pi'(a|s) - \pi(a|s)| \underbrace{\|P(\cdot|s, a)\|_1}_{=1} \right] \\ &= \mathbb{E}_{s \sim d_{\pi'}}[\|\pi'(\cdot|s) - \pi(\cdot|s)\|_1], \end{aligned} \quad (55)$$

which gives the claim. $\square$

For our mixture:

$$\|\pi_{\text{new}}(\cdot|s) - \pi(\cdot|s)\|_1 = \alpha\|\pi_{\text{cand}}(\cdot|s) - \pi(\cdot|s)\|_1 \leq 2\alpha, \quad (56)$$

since $\|p - q\|_1 \leq 2$ for any two probability distributions p, q on the same action space.

Thus:

$$\|d_{\pi_{\text{new}}} - d_{\pi}\|_1 \leq \frac{2\gamma\alpha}{1-\gamma}. \quad (57)$$

Combining the bounds. Write:

$$\mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)] = \mathbb{E}_{s \sim d_{\pi}}[\delta(s)] + (\mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)] - \mathbb{E}_{s \sim d_{\pi}}[\delta(s)]). \quad (58)$$

The second term is bounded by:

$$|\mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)] - \mathbb{E}_{s \sim d_{\pi}}[\delta(s)]| \leq \|d_{\pi_{\text{new}}} - d_{\pi}\|_1 \cdot \sup_s |\delta(s)| \leq \frac{2\gamma\alpha}{1-\gamma} \varepsilon_{\text{CPI}}. \quad (59)$$

Therefore:

$$\eta(\pi_{\text{new}}) = \eta(\pi) + \frac{\alpha}{1-\gamma} \mathbb{E}_{s \sim d_{\pi_{\text{new}}}}[\delta(s)] \quad (60)$$

$$\geq \eta(\pi) + \frac{\alpha}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}}[\delta(s)] - \frac{\alpha}{1-\gamma} \cdot \frac{2\gamma\alpha}{1-\gamma} \varepsilon_{\text{CPI}} \quad (61)$$

$$= L_{\pi}(\pi_{\text{new}}) - \frac{2\gamma\varepsilon_{\text{CPI}}}{(1-\gamma)^2} \alpha^2, \quad (62)$$

where in the last step we used $L_{\pi}(\pi_{\text{new}}) = \eta(\pi) + \frac{1}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}, a \sim \pi_{\text{new}}}[A^{\pi}(s, a)] = \eta(\pi) + \frac{\alpha}{1-\gamma} \mathbb{E}_{s \sim d_{\pi}}[\delta(s)]$.

This establishes (29).

Corollary K.2 (CPI with a centering defect). *Use the setup of Theorem 6.5, but replace the exact-centering condition (18) by*

$$\varepsilon_{\text{cent}} := \sup_{s \in \mathcal{R}_{\pi}} \left| \mathbb{E}_{a \sim \pi(\cdot|s)}[\hat{A}(s, a)] \right| < \infty.$$

Then

$$\eta(\pi_{\text{new}}) \geq \hat{L}_{\pi}(\pi_{\text{new}}) - \frac{(1-\alpha)\varepsilon_{\text{cent}} + \alpha\varepsilon_{A, \text{cand}}}{1-\gamma} - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1-\gamma)^2} \alpha^2. \quad (63)$$

Proof. The CPI shell gives

$$\eta(\pi_{\text{new}}) \geq L_{\pi}(\pi_{\text{new}}) - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1-\gamma)^2} \alpha^2.$$

For each $s \in \mathcal{R}_{\pi}$,

$$\begin{aligned} & \mathbb{E}_{a \sim \pi_{\text{new}}(\cdot|s)}[A^{\pi}(s, a) - \hat{A}(s, a)] \\ &= (1-\alpha)\mathbb{E}_{a \sim \pi(\cdot|s)}[A^{\pi}(s, a) - \hat{A}(s, a)] + \alpha\mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)}[A^{\pi}(s, a) - \hat{A}(s, a)]. \end{aligned}$$

Since $\mathbb{E}_{a \sim \pi}[A^{\pi}(s, a)] = 0$, the first term is bounded below by $-(1-\alpha)\varepsilon_{\text{cent}}$. By the definition of $\varepsilon_{A, \text{cand}}$, the second term is bounded below by $-\alpha\varepsilon_{A, \text{cand}}$. Therefore

$$L_{\pi}(\pi_{\text{new}}) \geq \hat{L}_{\pi}(\pi_{\text{new}}) - \frac{(1-\alpha)\varepsilon_{\text{cent}} + \alpha\varepsilon_{A, \text{cand}}}{1-\gamma}.$$

Combining the two displays proves (63). $\square$

K.1 PROOF OF COROLLARY 6.2

Proof. Because $\mathcal{T}_K^\pi$ is a γ^K -contraction with fixed point V^π ,

$$\begin{aligned} \|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} &\leq \|U_n - \mathcal{T}_K^\pi U_n\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} + \|\mathcal{T}_K^\pi U_n - \mathcal{T}_K^\pi V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \\ &\leq \varepsilon_{\text{res}, n} + \gamma^K \|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}. \end{aligned}$$

Rearranging gives (10). Lemma L.2 applied to the exact one-step centered estimator $\hat{A}_{U_n}$ then yields (11); by definition, $\varepsilon_{A, \text{cand}} \leq \varepsilon_A$. Substituting that bound into Theorem 6.5 gives (12). $\square$

Remark K.3 (Residual bridge via proxy quantities). The held-out U_n Bellman residual is a practical proxy for $\varepsilon_{\text{res}}^*$, but the global sup-norm quantity below is not directly measurable from finite data; it upper-bounds $\varepsilon_{\text{res}}^*$ up to the same truncation gap:

$$\varepsilon_{\text{res}}^* \leq \|U_n - \mathcal{T}_K^\pi U_n\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} + (1 + \gamma^K) \|U_n - U_*\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}. \quad (64)$$

This follows by adding and subtracting U_n inside $U_* - \mathcal{T}_K^\pi U_*$ and using K -step Bellman contraction on the final term. Combining (64) with Proposition A.2 gives

$$\varepsilon_{\text{res}}^* \leq \|U_n - \mathcal{T}_K^\pi U_n\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} + (1 + \gamma^K) L_V \frac{L_z^n}{1 - L_z} C_z. \quad (65)$$

L ADVANTAGE ESTIMATION AND BOUNDS

$$\hat{Q}_V(s, a) = \mathbb{E}_{s'}[r(s, a, s') + \gamma V(s')], \quad (66)$$

$$\hat{A}_V(s, a) = \hat{Q}_V(s, a) - \mathbb{E}_{b \sim \pi(\cdot|s)}[\hat{Q}_V(s, b)]. \quad (67)$$

Remark L.1 (Centering and baseline computation). The linear-in- α $\varepsilon_{A, \text{cand}}$ term in Theorem 6.5 relies on the *statewise centering* property

$$\mathbb{E}_{a \sim \pi(\cdot|s)}[\hat{A}(s, a)] = 0 \quad \text{for all } s.$$

This holds identically when $\hat{A}$ is defined as in (67) with the baseline computed exactly as $B_V(s) = \mathbb{E}_{b \sim \pi(\cdot|s)}[\hat{Q}_V(s, b)]$ (e.g., by summation over a finite action space).

If the baseline is approximated, centering may fail. We quantify this by the *centering defect*

$$\varepsilon_{\text{cent}} := \sup_{s \in \mathcal{R}_\pi} \left| \mathbb{E}_{a \sim \pi(\cdot|s)}[\hat{A}(s, a)] \right|.$$

Theorem 6.5 corresponds to the ideal case $\varepsilon_{\text{cent}} = 0$.

Lemma L.2 (Value to one-step advantage error). *Let $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ be the advantage-evaluation closure from Definition I.2, and let $\delta_V := \|V - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}$. Then for all (s, a) with $s \in \mathcal{R}_\pi$ and $a \in \text{supp}(\pi(\cdot|s)) \cup \text{supp}(\pi_{\text{cand}}(\cdot|s))$,*

$$|\hat{A}_V(s, a) - A^\pi(s, a)| \leq 2\gamma \delta_V.$$

The factor 2γ is tight worst-case (error appears in both Q and baseline). In practice, errors in $\hat{Q}_V(s, a)$ and the baseline $\mathbb{E}_\pi[\hat{Q}_V]$ often have the same sign and partially cancel.

Remark L.3 (Extensions to K -step and GAE-style advantages). Lemma L.2 is stated for the *one-step* estimator (67), hence the factor 2γ . If instead one defines a K -step action-value estimate

$$\hat{Q}_V^{(K)}(s, a) := \mathbb{E} \left[\sum_{t=0}^{K-1} \gamma^t r_t + \gamma^K V(s_K) \mid s_0 = s, a_0 = a, \pi \right],$$

and the corresponding centered advantage $\hat{A}_V^{(K)}(s, a) := \hat{Q}_V^{(K)}(s, a) - \mathbb{E}_{b \sim \pi(\cdot|s)}[\hat{Q}_V^{(K)}(s, b)]$ (with the baseline computed exactly), then the same argument yields

$$|\hat{A}_V^{(K)}(s, a) - A^\pi(s, a)| \leq 2\gamma^K \delta_V.$$

In practice we may use λ -return / GAE-style estimators (weighted mixtures of multi-step returns), which admit analogous bounds with a coefficient depending on (γ, λ) ; throughout, we absorb such estimator-dependent constants into the uniform advantage error parameter ε_A in Assumption 6.4.

Corollary L.4 (Advantage error in terms of theorem parameters). *Under Theorem 6.1 and the exact-centering conditions of Lemma L.2,*

$$\varepsilon_A \leq 2\gamma B_{n,m}. \quad (68)$$

Under the additional contractive and Lipschitz assumptions, taking $m \rightarrow \infty$ yields

$$\varepsilon_A \leq \underbrace{\frac{2\gamma \varepsilon_{\text{res}}^*}{1 - \gamma^K}}_{\text{architectural}} + \underbrace{2\gamma L_V \frac{L_z^n}{1 - L_z}}_{\text{unrolling}} C_z. \quad (69)$$

Here ε_A bounds the advantage error on $\mathcal{R}_\pi$ for actions $a \in \text{supp}(\pi(\cdot|s)) \cup \text{supp}(\pi_{\text{cand}}(\cdot|s))$, and the intermediate value-error quantity in Lemma L.2 is taken over the advantage-evaluation closure $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}}$ (Definition I.2).

M PROOFS

M.1 LEMMA M.1 (INTERIOR PROJECTION FIXED POINTS)

Lemma M.1 (Interior projection fixed points). *Let $\Pi_R : \mathbb{R}^d \rightarrow \mathbb{R}^d$ be the Euclidean projection onto the closed ball $\{z : \|z\|_2 \leq R\}$ as in (23). Suppose $z^* = \Pi_R(u)$ for some $u \in \mathbb{R}^d$. If $\|z^*\|_2 < R$, then $z^* = u$ (i.e., the projection is inactive at u). Consequently, if z^* is a fixed point of $f_\theta(z, y, x) = \Pi_R(f_\theta(z, y, x))$ and $\|z^*\|_2 < R$, then $f_\theta(z^*, y, x) = z^*$ and Π_R is inactive at $f_\theta(z^*, y, x)$.*

Proof. By (23), if $\|u\|_2 \leq R$ then $\Pi_R(u) = u$. Since $z^* = \Pi_R(u)$ lies in the open ball $\|z^*\|_2 < R$, we have $z^* \neq R \frac{u}{\|u\|_2}$, so the second case of (23) cannot apply; hence $\|u\|_2 \leq R$ and $\Pi_R(u) = u = z^*$.

For the consequence: if $z^* = \tilde{f}_\theta(z^*, y, x) = \Pi_R(f_\theta(z^*, y, x))$ and $\|z^*\|_2 < R$, applying the first statement with $u = f_\theta(z^*, y, x)$ gives $f_\theta(z^*, y, x) = z^*$. $\square$

M.2 PROOF OF PROPOSITION A.2

Proof. Fix (x, y) . By Assumption 5.1, $T_{x,y}(z) = \tilde{f}_\theta(z, y, x)$ maps $\mathcal{Z}_{\text{inv}}$ to itself (forward-invariance) and is a contraction with modulus $L_z < 1$. Since $\mathcal{Z}_{\text{inv}}$ is a closed subset of the complete metric space $(\mathcal{Z}, \|\cdot\|) = (\mathbb{R}^d, \|\cdot\|_2)$, it is itself complete.

Banach's theorem yields unique fixed point $z^*(x, y) \in \mathcal{Z}_{\text{inv}}$ and linear convergence (26).

For the finite-unrolling bound with $z^{(0)} = z_{\text{init}}(x, y)$:

$$z^* - z^{(n)} = \sum_{k=n}^{\infty} (z^{(k+1)} - z^{(k)}).$$

Contraction gives $\|z^{(k+1)} - z^{(k)}\| \leq L_z^k \|z^{(1)} - z^{(0)}\|$, so:

$$\|z^* - z^{(n)}\| \leq \|z^{(1)} - z^{(0)}\| \sum_{k=n}^{\infty} L_z^k = \frac{L_z^n}{1 - L_z} \|z^{(1)} - z^{(0)}\|.$$

By definition (25) (using the projected update $\tilde{f}_\theta = \Pi_R \circ f_\theta$), we have $\|z^{(1)} - z^{(0)}\| = \|\tilde{f}_\theta(z_{\text{init}}, y, x) - z_{\text{init}}\| \leq C_z$. $\square$

M.3 PROOF OF LEMMA B.7

Proof. For $T_{x,y_t}^{\circ n}$ a κ_n -contraction (where $\kappa_n = L_z^n$):

$$\|z_{t+1} - z^*(x, y_t)\| = \|T_{x,y_t}^{\circ n}(z_t) - T_{x,y_t}^{\circ n}(z^*(x, y_t))\| \leq \kappa_n e_t.$$

By Assumption B.3:

$$\|z^*(x, y_t) - z^*(x, y_{t+1})\| \leq L_{z^*} \|y_{t+1} - y_t\|_{\mathcal{E}} \leq L_{z^*} \Delta y_{\text{max}}.$$

Thus:

$$e_{t+1} = \|z_{t+1} - z^*(x, y_{t+1})\| \leq \|z_{t+1} - z^*(x, y_t)\| + \|z^*(x, y_t) - z^*(x, y_{t+1})\| \leq \kappa_n e_t + L_{z^*} \Delta y_{\max}.$$

Unrolling this recurrence:

$$e_t \leq \kappa_n^t e_0 + L_{z^*} \Delta y_{\max} \sum_{j=0}^{t-1} \kappa_n^j = \kappa_n^t e_0 + \frac{L_{z^*} \Delta y_{\max} (1 - \kappa_n^t)}{1 - \kappa_n}.$$

Using $e_0 \leq E_0$ and noting that this expression is maximized either at $t = 0$ (giving E_0) or as $t \rightarrow \infty$ (giving $\frac{L_{z^*} \Delta y_{\max}}{1 - \kappa_n}$), we have:

$$\sup_t e_t \leq \max \left(E_0, \frac{L_{z^*} \Delta y_{\max}}{1 - \kappa_n} \right) \leq E_0 + \frac{L_{z^*} \Delta y_{\max}}{1 - \kappa_n}.$$

For the used latent:

$$\|z_t^{(n)} - z^*(x, y_t)\| = \|T_{x, y_t}^{\circ n}(z_t) - z^*(x, y_t)\| = \|T_{x, y_t}^{\circ n}(z_t) - T_{x, y_t}^{\circ n}(z^*(x, y_t))\| \leq \kappa_n e_t,$$

where the second equality uses that $z^*(x, y_t)$ is the fixed point of T_{x, y_t} , hence also of $T_{x, y_t}^{\circ n}$, and the inequality follows from κ_n -contractivity of the n -fold composition. Thus:

$$\sup_t \|z_t^{(n)} - z^*(x, y_t)\| \leq \kappa_n \sup_t e_t \leq \kappa_n \max \left(E_0, \frac{L_{z^*} \Delta y_{\max}}{1 - \kappa_n} \right) = C_{\text{drift}}(n),$$

and the final “ $\leq$ ” in (42) gives the simpler additive upper bound. $\square$

M.4 PROOF OF THEOREM 6.1

Proof. Write $\|\cdot\|_\infty$ for the ordinary sup-norm on $\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} \cup \{s_{\text{abs}}\}$. The measurability, boundedness, and almost-sure forward-invariance assumptions ensure that $\mathcal{T}_K^\pi$ maps the relevant bounded functions back to this space and is a γ^K -contraction in this norm. Since $V^\pi = \mathcal{T}_K^\pi V^\pi$,

$$\begin{aligned} \|U_m - V^\pi\|_\infty &\leq \|U_m - \mathcal{T}_K^\pi U_m\|_\infty + \|\mathcal{T}_K^\pi U_m - \mathcal{T}_K^\pi V^\pi\|_\infty \\ &\leq \|U_m - \mathcal{T}_K^\pi U_m\|_\infty + \gamma^K \|U_m - V^\pi\|_\infty. \end{aligned}$$

Rearranging and then applying the triangle inequality to $U_n - U_m + U_m - V^\pi$ proves (2).

For every nonabsorbing s , Lipschitzness of the scalar head and the telescoping latent path give

$$\begin{aligned} |U_n(s) - U_m(s)| &\leq L_V \|z^{(n)}(s) - z^{(m)}(s)\| \\ &\leq L_V \sum_{j=n}^{m-1} \|z^{(j+1)}(s) - z^{(j)}(s)\|. \end{aligned}$$

The left side is zero at s_{abs} because its value is depth independent. Taking the supremum and using $\sup_s \sum_j g_j(s) \leq \sum_j \sup_s g_j(s)$ proves (3). The latent norm is the same norm fixed in the theorem, so each summand has latent units and multiplication by L_V gives value units.

Under Assumption 5.1,

$$\|z^{(j+1)}(s) - z^{(j)}(s)\| \leq L_z^j \|z^{(1)}(s) - z^{(0)}(s)\| \leq L_z^j C_z.$$

Summing from $j = n$ through $m - 1$ yields (5). Proposition A.2 and the Lipschitz head give $\|U_m - U_*\|_\infty \rightarrow 0$. Moreover,

$$\|U_m - \mathcal{T}_K^\pi U_m\|_\infty - \|U_* - \mathcal{T}_K^\pi U_*\|_\infty \leq (1 + \gamma^K) \|U_m - U_*\|_\infty \rightarrow 0.$$

Taking $m \rightarrow \infty$ therefore proves (6). $\square$

Remark M.2 (Finite- R specialization of the truncation constant). Under the Appendix N initialization condition $\|z_{\text{init}}(x, y)\|_2 \leq R$ with $\mathcal{Z}_{\text{inv}} = \{z : \|z\|_2 \leq R\}$, the projection Π_R ensures $\|\tilde{f}_\theta(z_{\text{init}}(x, y), y, x)\|_2 \leq R$ for all (x, y) . Therefore

$$C_z = \sup_{(x, y)} \|\tilde{f}_\theta(z_{\text{init}}(x, y), y, x) - z_{\text{init}}(x, y)\|_2 \leq 2R.$$

The contractive specialization of Theorem 6.1 therefore gives

$$\|U_n - V^\pi\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}} \leq \frac{\varepsilon_{\text{res}}^*}{1 - \gamma^K} + 2L_V \frac{L_z^n}{1 - L_z} R.$$

This makes R explicit only through the finite-unrolling term; it does not identify the dependence of $\varepsilon_{\text{res}}^*$ on R .

M.5 PROJECTION-ACTIVE THEORY

Projection-aware CPI penalty (full statement). In the projection-active regime, assuming $L_z^{\text{sat}}(R) < 1$ in Proposition M.3, combining (7), (8), and $C_z \leq 2R$ yields a projection-aware version of (22):

$$\begin{aligned} \eta(\pi_{\text{new}}) &\geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{2\alpha\gamma}{1 - \gamma} \left(\frac{\varepsilon_0 + (1 + \gamma^K)L_V(R)R}{1 - \gamma^K} + 2L_V(R) \frac{(L_z^{\text{sat}}(R))^n}{1 - L_z^{\text{sat}}(R)} R \right) \\ &\quad - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2. \end{aligned} \quad (70)$$

Projection affects both the architectural and truncation terms; the finite- R sweep (Appendix D.5) shows this is a regime diagnostic, not a monotone tuning rule for R .

Proposition M.3 (Projection-induced contraction on a saturated annulus). *Let $B_R := \{z : \|z\|_2 \leq R\}$ and define the pre-projection local Lipschitz constant*

$$L_z^{\text{pre}}(R) := \sup_{(x, y)} \sup_{\substack{z, z' \in B_R \\ z \neq z'}} \frac{\|f_\theta(z, y, x) - f_\theta(z', y, x)\|_2}{\|z - z'\|_2}.$$

Assume there exists $\rho_R > R$ such that

$$\|f_\theta(z, y, x)\|_2 \geq \rho_R \quad \text{for all relevant } (x, y) \text{ and } z \in B_R.$$

Then the projected map $\tilde{f}_\theta = \Pi_R \circ f_\theta$ is Lipschitz on B_R with modulus

$$L_z^{\text{sat}}(R) \leq \frac{R}{\rho_R} L_z^{\text{pre}}(R).$$

In particular, if $(R/\rho_R) L_z^{\text{pre}}(R) < 1$, then the contraction part of Assumption 5.1 holds on B_R with $L_z \leq L_z^{\text{sat}}(R)$; if moreover $z_{\text{init}}(x, y) \in B_R$ for all (x, y) , then Assumption 5.1 holds in full on B_R .

Proof. For any u, v with $\|u\|_2, \|v\|_2 \geq \rho_R > R$,

$$\|\Pi_R(u) - \Pi_R(v)\|_2 = R \left\| \frac{u}{\|u\|_2} - \frac{v}{\|v\|_2} \right\|_2.$$

Using

$$\left\| \frac{u}{\|u\|_2} - \frac{v}{\|v\|_2} \right\|_2^2 = \frac{\|u - v\|_2^2 - (\|u\|_2 - \|v\|_2)^2}{\|u\|_2 \|v\|_2} \leq \frac{\|u - v\|_2^2}{\rho_R^2},$$

we obtain

$$\|\Pi_R(u) - \Pi_R(v)\|_2 \leq \frac{R}{\rho_R} \|u - v\|_2.$$

Applying this with $u = f_\theta(z, y, x)$ and $v = f_\theta(z', y, x)$ for $z, z' \in B_R$ gives

$$\|\tilde{f}_\theta(z, y, x) - \tilde{f}_\theta(z', y, x)\|_2 \leq \frac{R}{\rho_R} L_z^{\text{pre}}(R) \|z - z'\|_2. \quad \square$$

Proposition M.4 (Projection-aware architectural residual via a plan-independent latent anchor). *Fix any anchor map $z_{\text{anc}} : \mathcal{X} \rightarrow B_R$ that depends only on the instance x , and define*

$$B_{\text{anc}}(x, y) := V_{\psi}(z_{\text{anc}}(x), x), \quad B_{\text{anc}}(s_{\text{abs}}) := -C_{\text{max}},$$

$$\varepsilon_{\text{anc}} := \|B_{\text{anc}} - \mathcal{T}_K^{\pi} B_{\text{anc}}\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}.$$

Let

$$D_{\text{anc}}(R) := \sup_{(x, y)} \|z^*(x, y) - z_{\text{anc}}(x)\|_2, \quad L_V(R) := \sup_x \sup_{\substack{z, z' \in B_R \\ z \neq z'}} \frac{|V_{\psi}(z, x) - V_{\psi}(z', x)|}{\|z - z'\|_2}.$$

Then

$$\varepsilon_{\text{res}}^* \leq \varepsilon_{\text{anc}} + (1 + \gamma^K) L_V(R) D_{\text{anc}}(R). \quad (71)$$

In the special case $z_{\text{anc}}(x) \equiv 0$,

$$B_0(x, y) := V_{\psi}(0, x), \quad B_0(s_{\text{abs}}) := -C_{\text{max}},$$

$$\varepsilon_{\text{res}}^* \leq \varepsilon_0 + (1 + \gamma^K) L_V(R) R,$$

$$\varepsilon_0 := \|B_0 - \mathcal{T}_K^{\pi} B_0\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}.$$

Proof. Add and subtract B_{anc} :

$$U_* - \mathcal{T}_K^{\pi} U_* = (B_{\text{anc}} - \mathcal{T}_K^{\pi} B_{\text{anc}}) + (U_* - B_{\text{anc}}) + (\mathcal{T}_K^{\pi} B_{\text{anc}} - \mathcal{T}_K^{\pi} U_*).$$

Taking $\|\cdot\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}$ and using the γ^K -contraction of $\mathcal{T}_K^{\pi}$ on the difference of its two images gives

$$\varepsilon_{\text{res}}^* \leq \varepsilon_{\text{anc}} + (1 + \gamma^K) \|U_* - B_{\text{anc}}\|_{\infty, \pi, \pi_{\text{cand}}}^{\text{adv}}.$$

Since both $z^*(x, y)$ and $z_{\text{anc}}(x)$ lie in B_R ,

$$|U_*(x, y) - B_{\text{anc}}(x, y)| = |V_{\psi}(z^*(x, y), x) - V_{\psi}(z_{\text{anc}}(x), x)|$$

$$\leq L_V(R) \|z^*(x, y) - z_{\text{anc}}(x)\|_2 \leq L_V(R) D_{\text{anc}}(R).$$

At the absorbing boundary, $U_*(s_{\text{abs}}) = B_{\text{anc}}(s_{\text{abs}}) = -C_{\text{max}}$, so the difference is zero. The special case $z_{\text{anc}}(x) \equiv 0$ uses $D_{\text{anc}}(R) \leq R$. $\square$

M.6 PROOF OF LEMMA L.2

Proof. For any (s, a) ,

$$\widehat{Q}_V(s, a) - Q^{\pi}(s, a) = \gamma \mathbb{E}_{s' \sim P(\cdot | s, a)} [V(s') - V^{\pi}(s')].$$

If $s \in \mathcal{R}_{\pi}$ and $a \in \text{supp}(\pi(\cdot | s)) \cup \text{supp}(\pi_{\text{cand}}(\cdot | s))$, Definition I.2 gives $P(\mathcal{R}_{\pi, \pi_{\text{cand}}}^{\text{adv}} | s, a) = 1$. Thus $|V(s') - V^{\pi}(s')| \leq \delta_V$ for transition-almost every s' , and hence $|\widehat{Q}_V(s, a) - Q^{\pi}(s, a)| \leq \gamma \delta_V$.

Similarly,

$$|\mathbb{E}_{b \sim \pi(\cdot | s)} \widehat{Q}_V(s, b) - \mathbb{E}_{b \sim \pi(\cdot | s)} Q^{\pi}(s, b)| \leq \gamma \delta_V,$$

since $b \sim \pi(\cdot | s)$ and the same almost-sure closure condition applies. Subtracting these two bounds yields $|\widehat{A}_V(s, a) - A^{\pi}(s, a)| \leq 2\gamma \delta_V$. $\square$

M.7 PROOF OF THEOREM 6.5

Proof. The CPI bound (29) (derived in Appendix K) gives:

$$\eta(\pi_{\text{new}}) \geq L_{\pi}(\pi_{\text{new}}) - \frac{2\varepsilon_{\text{CPI}}\gamma}{(1 - \gamma)^2} \alpha^2.$$

Now relate the true surrogate $L_{\pi}(\pi_{\text{new}})$ to the estimated surrogate $\widehat{L}_{\pi}(\pi_{\text{new}})$:

$$L_{\pi}(\pi_{\text{new}}) - \widehat{L}_{\pi}(\pi_{\text{new}}) = \frac{1}{1 - \gamma} \mathbb{E}_{s \sim d_{\pi}, a \sim \pi_{\text{new}}(\cdot | s)} [A^{\pi}(s, a) - \widehat{A}(s, a)].$$

For the mixture policy $\pi_{\text{new}} = (1 - \alpha)\pi + \alpha\pi_{\text{cand}}$:

$$\begin{aligned}\mathbb{E}_{a \sim \pi_{\text{new}}(\cdot|s)}[A^\pi(s, a) - \hat{A}(s, a)] &= (1 - \alpha)\mathbb{E}_{a \sim \pi(\cdot|s)}[A^\pi(s, a) - \hat{A}(s, a)] \\ &\quad + \alpha\mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)}[A^\pi(s, a) - \hat{A}(s, a)].\end{aligned}$$

By definition of advantage, $\mathbb{E}_{a \sim \pi(\cdot|s)}[A^\pi(s, a)] = 0$. By the centering assumption (18), $\mathbb{E}_{a \sim \pi(\cdot|s)}[\hat{A}(s, a)] = 0$. Hence the $(1 - \alpha)$ term vanishes, and we obtain

$$\left| \mathbb{E}_{a \sim \pi_{\text{new}}(\cdot|s)}[A^\pi(s, a) - \hat{A}(s, a)] \right| = \alpha \left| \mathbb{E}_{a \sim \pi_{\text{cand}}(\cdot|s)}[A^\pi(s, a) - \hat{A}(s, a)] \right| \leq \alpha \varepsilon_{A, \text{cand}}.$$

Why this is a tightening: Note that by Jensen’s inequality and Assumption 6.4,

$$\varepsilon_{A, \text{cand}} = \sup_s \left| \mathbb{E}_{a \sim \pi_{\text{cand}}}[\hat{A}(s, a) - A^\pi(s, a)] \right| \leq \sup_s \mathbb{E}_{a \sim \pi_{\text{cand}}} [|\hat{A}(s, a) - A^\pi(s, a)|] \leq \varepsilon_A.$$

Thus the evaluation error contribution $\alpha \varepsilon_{A, \text{cand}} / (1 - \gamma)$ is no larger than the naive bound $\varepsilon_A / (1 - \gamma)$, and is smaller whenever $\alpha < 1$ or $\varepsilon_{A, \text{cand}} < \varepsilon_A$.

Taking expectation over $s \sim d_\pi$ yields

$$L_\pi(\pi_{\text{new}}) \geq \hat{L}_\pi(\pi_{\text{new}}) - \frac{\alpha}{1 - \gamma} \varepsilon_{A, \text{cand}}.$$

Combining with the CPI bound gives (20). $\square$

M.8 PROOF OF COROLLARY B.1

Proof. The clock-complete state, frozen shared map h_n , and specified initial law define one augmented MDP for both action kernels. On its bounded measurable functions, $\bar{\mathcal{T}}_K^{\bar{\pi}}$ is a γ^K -contraction in the augmented closure norm. In this proof, $\|\cdot\|_\infty^{\text{adv}}$ abbreviates $\|\cdot\|_{\infty, \bar{\pi}, \bar{\pi}_{\text{cand}}}^{\text{adv}}$. Thus

$$\begin{aligned}\|\tilde{U}_n - V^{\bar{\pi}}\|_\infty^{\text{adv}} &\leq \|\tilde{U}_n - \bar{\mathcal{T}}_K^{\bar{\pi}} \tilde{U}_n\|_\infty^{\text{adv}} + \|\bar{\mathcal{T}}_K^{\bar{\pi}} \tilde{U}_n - \bar{\mathcal{T}}_K^{\bar{\pi}} V^{\bar{\pi}}\|_\infty^{\text{adv}} \\ &\leq \bar{\varepsilon}_{\text{res}, n} + \gamma^K \|\tilde{U}_n - V^{\bar{\pi}}\|_\infty^{\text{adv}}.\end{aligned}$$

Rearrangement proves (37). Lemma L.2, applied on the augmented closure with exact centering conditional on $\bar{s}$, gives (38). The candidate-policy bias is no larger than this uniform advantage error. Under the exact pointwise mixture, substituting it into Theorem 6.5 proves (39). $\square$

For completeness, the optional slow-drift comparisons follow directly. Lemma B.7 and the Lipschitz head imply (44); adding and subtracting $\tilde{U}_*$ and applying the residual bound (1) proves (45). Finally, the reverse triangle inequality and Bellman contraction give

$$\begin{aligned}|\bar{\varepsilon}_{\text{res}, n} - \bar{\varepsilon}_{\text{res}, *}| &\leq \|(\tilde{U}_n - \tilde{U}_*) - (\bar{\mathcal{T}}_K^{\bar{\pi}} \tilde{U}_n - \bar{\mathcal{T}}_K^{\bar{\pi}} \tilde{U}_*)\|_\infty^{\text{adv}} \\ &\leq (1 + \gamma^K) \|\tilde{U}_n - \tilde{U}_*\|_\infty^{\text{adv}},\end{aligned}$$

which proves (46).

N IMPLEMENTATION DETAILS

N.1 ALGORITHM PSEUDOCODE

Algorithm 1 presents the full UPI-TRM update step in persistent- z form. For the episodic- z variant, skip line 10 (“Carry latent”) and reinitialize $z_k \leftarrow z_{\text{init}}(x, y_k)$ at the start of each loop iteration.

The frozen CPI comparison begins only after the value update and uses the same fixed recurrent map h_n for the current and candidate action kernels through the mixture step. Any subsequent recurrent-parameter update starts a new snapshot and is outside Theorem 6.5. The retained training repository is unavailable, so whether the headline implementation enforces this separation remains unaudited.

Algorithm 2 UPI-TRM: single interaction + update step

```

1: Input:  $x, y_0, z_0, \gamma, n, K, \alpha, \tau, R$ 
2: Initialize  $\tilde{s}_0 \leftarrow (x, y_0, z_0)$ ;  $done \leftarrow false$ 
3: for  $k = 0$  to  $K - 1$  do
4:   if  $done$  then
5:      $r_k \leftarrow (\gamma - 1)C_{\max}$ ;  $\tilde{s}_{k+1} \leftarrow s_{\text{abs}}$  // Absorbing self-loop
6:   else
7:      $z_k^{(n)} \leftarrow (\Pi_R \circ f_\theta)^{\circ n}(z_k, y_k, x)$  // n inner steps
8:     Sample  $a_k \sim \pi_\phi(a \mid y_k, z_k^{(n)}, x)$ 
9:      $y_{k+1} \leftarrow \text{edit}(y_k, a_k; x)$ ; receive  $r_k$ 
10:     $z_{k+1} \leftarrow z_k^{(n)}$  // Carry latent (skip for episodic-z)
11:     $\tilde{s}_{k+1} \leftarrow (x, y_{k+1}, z_{k+1})$ 
12:    if terminated then
13:       $done \leftarrow true$ ;  $\tilde{s}_{k+1} \leftarrow s_{\text{abs}}$ 
14:    end if
15:  end if
16: end for
17: Compute target  $G^{(K)}$  via (24)
18: Update  $(\theta, \psi)$  by minimizing  $\mathcal{L}_{\text{val}}$ , then freeze the evaluator snapshot  $h_n$ 
19: Form advantages  $\hat{A}$  via centered estimator (see below)
20: Policy-head step with  $h_n$  frozen  $\rightarrow$  candidate  $\pi_{\text{cand}}$ 
21:  $\pi_{\text{new}} \leftarrow (1 - \alpha)\pi_\phi + \alpha\pi_{\text{cand}}$ 
22: Distill:  $\pi_\phi \leftarrow \arg \min_\phi \text{KL}(\pi_{\text{new}} \parallel \pi_\phi)$ 
23:  $\bar{\psi} \leftarrow \tau\bar{\psi} + (1 - \tau)\psi$  // Target network update

```

We maintain a replay buffer and sample mini-batches for value and policy updates. Latents in the buffer computed under older θ are stale, so we recompute $z^{(n)}$ for sampled states or use a slowly moving target encoder $f_{\bar{\theta}}$.

Compute and hardware. Each reported training or evaluation run uses a single device. The main training jobs and the trusted external PPO/A2C/DQN reruns, including the 5-step DQN variant, ran through GPU-enabled batch jobs; the evaluation scripts also support CPU fallback for analysis-only runs. Within the architecture-matched group, “matched” refers to the shared TRM backbone and checker-shaped reward, not to identical environment-interaction counts or constant-factor wall-clock cost. The nominal 20k outer-step protocol gives different interaction budgets across PPO, A2C, and DQN because each collects a different number of transitions per outer step; Appendix G.1 reports the missing-audit and artifact status, not validated equalized results. The baseline update rules also carry algorithm-specific overhead; for example, PPO and A2C evaluate policy and value heads in separate forwards through the TRM backbone during optimization, so wall-clock time per update can exceed UPI-TRM.

Advantage estimation. The theory uses the centered estimator (67); in practice GAE/ (λ) -return variants are also fine for variance reduction. The linear-in- α error term in Theorem 6.5 requires *exact* statewise centering $\mathbb{E}_{a \sim \pi}[\hat{A}(s, a)] = 0$. This holds whenever the estimator is constructed to be exactly centered. For Eq. (67), the baseline $B_V(s) = \mathbb{E}_{b \sim \pi(\cdot|s)}[\hat{Q}_V(s, b)]$ must be evaluated exactly; the finite discrete implementation uses explicit summation. With a learned value baseline (as in standard GAE), centering is only approximate and the deterministic bound degrades by the centering defect $\varepsilon_{\text{cent}}$; the naive $O(\varepsilon_A/(1 - \gamma))$ penalty is recovered only as a worst-case upper bound when $\varepsilon_{\text{cent}}$ and $\varepsilon_{A, \text{cand}}$ are both controlled by ε_A .

Operator-norm control and contraction-oriented details. We apply explicit operator-norm clamping on the inner $z \rightarrow z$ recursion (periodic rescaling of linear weights using power-iteration estimates), followed by per-layer output scaling or an explicit scalar factor targeting $L_z < 1$. This intervention does not certify the uniform global modulus in Assumption 5.1. We estimate only a local mini-batch proxy $\hat{L}_z$ by finite differences:

$$\hat{L}_z = \max_i \frac{\|f_\theta(z_i, y_i, x_i) - f_\theta(z'_i, y_i, x_i)\|}{\|z_i - z'_i\|}.$$

Enforcing forward-invariance. To ensure $z_{\text{init}}(x, y) \in \mathcal{Z}_{\text{inv}}$ and that iterates remain there:

- Use an initialization $z_{\text{init}}(x, y)$ with $\|z_{\text{init}}(x, y)\|_2 \leq R$, $\mathcal{Z}_{\text{inv}} = \{z : \|z\|_2 \leq R\}$. In the artifact implementation, the initial latent comes from learned latent seeds (optionally augmented by an initialization encoder) and, when $R > 0$, is projected into the same ball.
- After each update, apply Euclidean projection onto the ball: $z \leftarrow \Pi_R(z)$ with Π_R as in (23).
- Under the Euclidean norm this projection is 1-Lipschitz (non-expansive) and so preserves the contraction modulus.
- Equivalently, define $\tilde{f}_\theta := \Pi_R \circ f_\theta$ and unroll $\tilde{f}_\theta^n$: $z^{(t+1)} = \Pi_R(f_\theta(z^{(t)}, y, x))$ for $t = 0, \dots, n-1$.

For persistent- z we monitor the fresh-initialization discrepancy $\hat{C}_{\text{fresh}}(n) = \max_t \|z_t^{(n)} - z_{\text{init}}(x, y_t)\|$. It measures separation from a fresh initialization; it is not a fixed-point tracking bound and is not a substitute for $C_{\text{drift}}(n)$ or the augmented residual.

Distillation. When distilling π_{new} to π_ϕ , we minimize:

$$\mathcal{L}_{\text{distill}}(\phi) = \mathbb{E}_{\tilde{s} \sim \text{buffer}} [\text{KL}(\pi_{\text{new}}(\cdot | \tilde{s}) \| \pi_\phi(\cdot | \tilde{s}))].$$

This is a heuristic outside the theory. A finite-batch $\text{KL}(\pi_{\text{new}} \| \pi_\phi)$ should be reported when both policies are retained, but it would still not certify exact deployment. That diagnostic is not available for the headline run, so its distillation gap is unmeasured rather than zero.

O MAPPING TO ORIGINAL TRM ARCHITECTURE

We provide an explicit mapping between our formalism and the original TRM implementation (Jolicoeur-Martineau, 2025).

Our Formalism	TRM Implementation	Notes
$z \in \mathcal{Z}$	Hidden state	Latent reasoning state
$y \in \mathcal{Y}$	Output grid/sequence	Candidate solution
f_θ	Recurrent update	Single update step
n (inner steps)	Recursion depth	Fixed or adaptively chosen evaluation budget
$z_{\text{init}}(x, y)$	Initial hidden	Learned (projected when $R > 0$)
$\pi_\phi(a y, z, x)$	Task head (reinterpreted)	TRM’s task head produces the refined answer; we reframe the answer update as a discrete edit policy
$\text{edit}(y, a; x)$	Continuous refinement	TRM: continuous y -update via another forward pass; we discretize into edit actions
Episodic- z	Not in original	Our simplification
Persistent- z	Original TRM	z carried across edits
Supervised objective	Training signal	Replaced by RL reward
Adaptive halting	Early stopping	Can be related to variable n

The key differences:

1. **Training signal:** The original TRM is trained with a supervised loss; we replace this with checker-based rewards and a CPI-style policy-improvement loop.
2. **Episodic- z vs. persistent- z :** We analyze both; the original TRM is persistent- z .
3. **Halting:** The original TRM uses an adaptive halting mechanism; our analysis fixes n for tractability and treats n as an analysis parameter rather than a learned quantity.

P ETHICS AND BROADER IMPACT

Ethics statement. The paper studies checker-trained Tiny Recursive Models on synthetic Sudoku tasks. We use no personal data, no human subjects, no crowdsourcing, no deployed decision system. The main ethical risk is rhetorical: narrow verifier-feedback results on puzzle tasks can be overstated as evidence about general reasoning or high-stakes automation. We scope the empirical claim to

1998 controlled Sudoku, do not include the 9×9 feasibility sweep as a result, and state in the limitations
1999 paragraph that the bound does not certify improvement.
2000
2001 **Broader impact.** The methodological contribution is to separate internal evaluator depth from
2002 outer policy improvement, which may help analyses of verifier-feedback systems with bounded
2003 internal compute. The main risk is misinterpretation of puzzle-domain results as broader claims about
2004 reasoning capability; we therefore avoid deployment recommendations and keep the paper’s claims
2005 narrow.
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051